


# CS330: Introduction

# Who am I?

- Youngjin Kwon, Assistant Professor at KAIST
  - who has studied operating systems
  - who loves designing/building systems
  - who enjoys low-level details about computer system
- Office: E3-1, 4405
- Email: [yjkwon@cs.kaist.ac.kr](mailto:yjkwon@cs.kaist.ac.kr)

# Prerequisites

- Prerequisites
  - CS230 is strongly recommended, and CS211 is preferably recommended



# Prerequisites

- You are expected to have
  - Basic understanding of key OS concepts
    - process, system call, virtual memory, file
    - The schedule of programming project goes ahead of lecture
  - Entry level of C programming experience
    - e.g., pointers to structures/arrays
    - Linux programming environment (GNU C)
    - Systems calls, file I/O programming, signal handling

547 -

# Logistics

- Textbook
  - Operating Systems Principles & Practice


- Discussion board

<https://piazza.com/kaist.ac.kr/spring2019/cs330/home>

- Course webpage

<https://sites.google.com/view/cs330/>

<http://klms.kaist.ac.kr>



# Learning pyramid

- A various percentage of retention of learning with different ways of teaching methods
- Studied by National Training Laboratory

	% of retention
Lecturing (listening)	5%
Reading	10%
Demonstration	30%
Discussion	50%
Practice by doing	70%
Teaching others	90%

# Credit criteria

- Mid-term exam: 20%
- Hands-on labs: 40%
- Final exam: 30%
- Pintos exam: 10%

# Project: Pintos

- Educational operating system
  - Simplified OS with essential features
- We give you some basic building blocks
  - Four assignments, that build on each other
    - Threads, user programs, virtual memory
  - Work in **groups of 2**  
Pair-programming:  
[https://en.wikipedia.org/wiki/Pair\\_programming](https://en.wikipedia.org/wiki/Pair_programming)



- Project document

## Short Contents

1	<u>Introduction</u> . . . . .	1
2	<u>Project 1: Threads</u> . . . . .	9
3	Project 2: User Programs . . . . .	22
4	Project 3: Virtual Memory . . . . .	39
5	Project 4: File Systems . . . . .	50
A	Reference Guide . . . . .	58

# Project: Pintos

- How to download pinto code will be announce in the discussion board or course webpage
- Four implementation projects:
  - Threads (Thread support, synchronization)
  - User programs
  - Virtual memory
  - File systems
- Project Guide & Help Sessions by TA
  - Wednesday, 7pm – 10pm

# Project Sessions (by TAs)

- Project sessions (online help)
  - **Lab hours** for programming help
    - Programming help for 25 min per team
    - You must request by one day before the lab hours day, describing the problem you want to ask
      - Please make sure your question is specific and clear
        - » TAs cannot help you if your question is too general
        - » Good example: *In the test case A, I have a problem at the line #XX*
          - *with explanation of your design and where you changed*
        - » Bad example: *My Pintos is not working*
- Due date is **Sunday midnight in general**
  - A late penalty is 20% off per day (up to 3 days)
  - 5 tokens for an entire semester
    - You can use a token for penalty waiver
  - Once your team submits a project after using all the penalty waivers, you will still get 35% of your score (**do not give up!**)

# Course Homepage

- Three assignments spread over quarter
  - Practice for final
  - Done **individually**

# Main Points (for today)

- Operating system definition
  - Software to manage a computer's resources for its users and applications
- OS challenges
  - Reliability, security, responsiveness, portability, ...
- How to understand OS?

# What's going on?

```
#include <stdio.h>
#include <string.h>
#include <sys/mman.h>
#include <sys/time.h>

int main(void)
{
    void *addr;
    struct timeval start, end, elap;

    addr = mmap(NULL, 1 << 20, PROT_READ | PROT_WRITE,
        MAP_ANONYMOUS | MAP_PRIVATE, -1, 0);

    gettimeofday(&start, NULL);
    memset(addr, 1, 1 << 20);
    gettimeofday(&end, NULL);

    timersub(&end, &start, &elap);

    printf("Time taken to memset1 %ld usec\n", elap.tv_usec);

    gettimeofday(&start, NULL);
    memset(addr, 2, 1 << 20);
    gettimeofday(&end, NULL);

    timersub(&end, &start, &elap);

    printf("Time taken to memset2 %ld usec\n", elap.tv_usec);

    munmap(addr, 1 << 20);

    return 0;
}
```

```
[ yjkwon@tigris02 ~] > gcc -o map map.c
[ yjkwon@tigris02 ~] > ./map
Time taken to memset1 1389 usec
Time taken to memset2 112 usec
[ yjkwon@tigris02 ~] > █
```

Lesson from computer architecture..  
CPU cache locality

# MAP size is way beyond CPU cache size!

```
#include <stdio.h>
#include <string.h>
#include <sys/mman.h>
#include <sys/time.h>

#define MAP_SIZE (1 << 30)

int main(void)
{
    void *addr;
    struct timeval start, end, elap;

    addr = mmap(NULL, MAP_SIZE, PROT_READ | PROT_WRITE,
        MAP_ANONYMOUS | MAP_PRIVATE, -1, 0);

    gettimeofday(&start, NULL);
    memset(addr, 1, MAP_SIZE);
    gettimeofday(&end, NULL);

    timersub(&end, &start, &elap);

    printf("Time taken to memset1 %0.2lf msec\n",
        (((double)elap.tv_sec * 1000000.0) + (double)elap.tv_usec) / 1000.0 );

    gettimeofday(&start, NULL);
    memset(addr, 2, MAP_SIZE);
    gettimeofday(&end, NULL);

    timersub(&end, &start, &elap);

    printf("Time taken to memset2 %0.2lf msec\n",
        (((double)elap.tv_sec * 1000000.0) + (double)elap.tv_usec) / 1000.0 );

    munmap(addr, MAP_SIZE);

    return 0;
}
```

```
[ yjkwon@tigris02 ~] > !gcc
[ yjkwon@tigris02 ~] > gcc -o map map.c
[ yjkwon@tigris02 ~] > ./map
Time taken to memset1 296.62 msec
Time taken to memset2 155.95 msec
[ yjkwon@tigris02 ~] > █
```

Why?


# Why is studying OS useful?

- “How operating system works” is essential for any students interested in build modern computer systems like web server, KV store etc
- ***OS as a reference material***
  - Software engineers use many of the same techniques and design patterns as those use in OS




# Main Role of OS

- To programs,
  - Providing application programming interface (API) to use hardware
  - Hide details of hardware



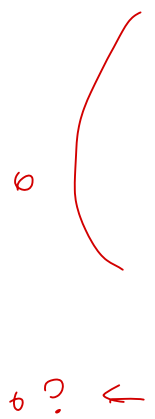
Operating system



# OS definition

- An Operating System is a layer of systems software that
  - ✱ – Has (privileged access) to the hardware directly;
  - Hide the hardware complexity;
  - Manages hardware on behalf of applications;
  - Additionally, ensures that applications are isolated and protected each other

# Question

- Which of the followings are likely components of an OS?
    - File editor
    - Browser
    - File system
    - Device D river
    - Cache memory
    - Process Scheduler
    - C standard library
- 

Design + Implementation

# Operating System Roles

- **Referee:** managing hardware and application activities
  - Resource allocation among users, applications
  - Isolation of different users, applications from each other
  - Communication between users, applications
- **Illusionist:** masking limitations of hardware
  - Each application appears to have the entire machine to itself
  - Infinite number of processors, (near) infinite amount of memory, reliable storage, reliable network transport
- **Glue:** providing a set of common, standard services to applications
  - Interface to standard C library, Graphic interfaces

# OS Challenges

Goal

→ Most dear & intuitive goal

## • Performance

Desktop prefers latency than throughput

metrics

### – Latency/response time

- How long does an operation take to complete?

### – Throughput → Server prefers

- How many operations can be done per unit of time?

### – Overhead

- How much extra work is done by the OS?

### – Fairness

- How equal is the performance received by different users?

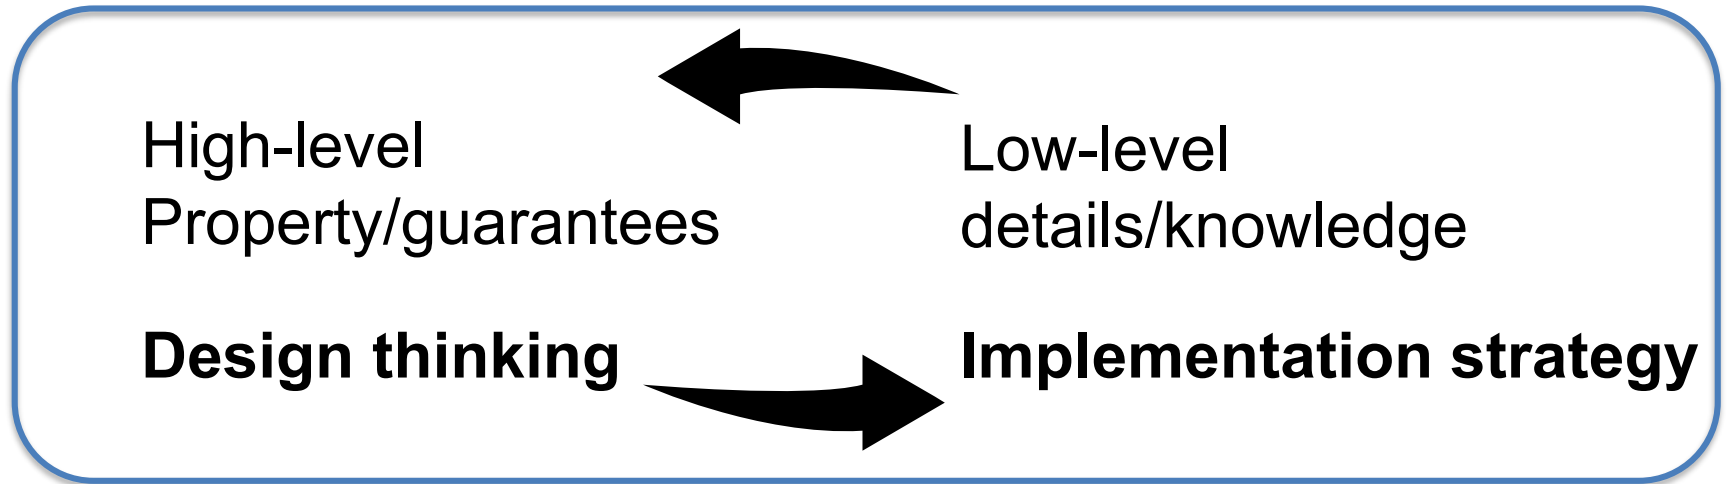
### – Predictability

- How consistent is the performance over time?

# OS Challenges

- **Security** - Write  
– Can the system be compromised by an attacker?  
*권한 없음.*
- **Privacy** - Read  
– Data is accessible only to authorized users
- **Reliability**  
– Does the system do what it was designed to do?
- **Availability**  
– What portion of the time is the system working?  
– Mean Time To Failure (MTTF), Mean Time to Repair

# How to understand OS?



*"Design Rationale"*

For example, **Protection**


- The goal of protection
- Protection model/API
- Efficient mechanisms

- How to use hardware mechanism?
- What part is done by hardware? to be done by software?
- How to separate software layers?

# OS in a bird's-eye view

OS guarantees:

- 1.
  - 2.
  - 3.
  - 4.
- You need to fill them at the end of the semester!



Operating system

