

The Programming Interface

Instructor: Youngjin Kwon

Let's think about interface

- OS designer
 - Know internals of how OS works
 - Know internals of how hardware works
- Application developer
 - Do not want to know internals of OS and hardware
- OS provides a simplistic view to use and manage OS internal resources by **APIs**
 - System calls

System call

What to consider designing APIs?

- What functionality should be in kernel?
 - What functionality should be in user library?
- How should OS be organized?

Some thoughts

- Let's focus on building APIs for process creation
- We can put the functionality of managing processes and IO in user-level
 - *What is a problem?*

Security ...?


Some thoughts (Cont'd)

- So, the functionality must be in the kernel
- Then, how does the APIs look like?
 - Designing good APIs is a hard problem
 - Our approach: Let's look at use cases and build APIs to reflect the use cases


Main Points

- Creating and managing processes
 - fork, exec, wait
- Performing I/O
 - open, read, write, close
- Monolithic kernel & Microkernel

Use case: Shell



- A shell is a job control system
 - Allows programmer to create and manage a set of programs to do some tasks
 - Windows, MacOS, Linux all have shells



- Example: to compile a C program

`$ cc -c sourcefile1.c` ➔ It requires create a process and write data to sourcefile1.o

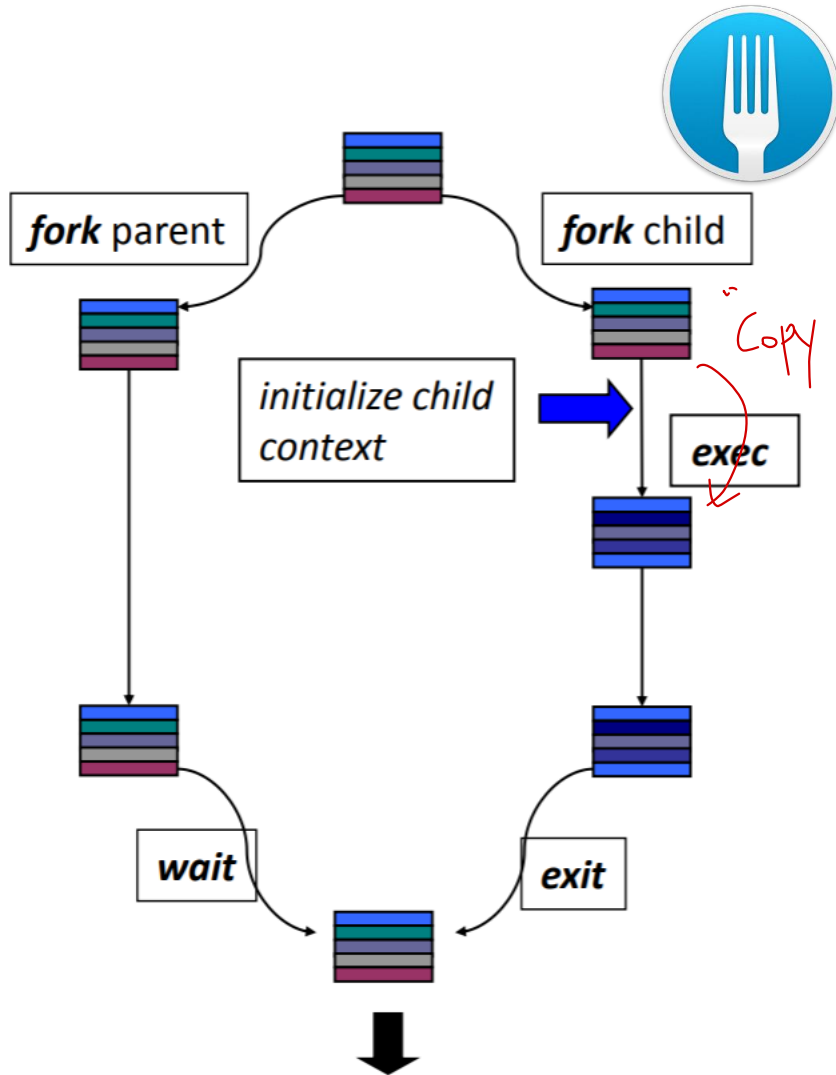
`$ cc -c sourcefile2.c`

`$ ln -o program sourcefile1.o sourcefile2.o`

Design questions

- If the shell runs at user-level, what system calls does it make to run each of the programs?
 - Ex: gcc, ln
- Why does process creation need system calls?
 - Why kernel-level? → for protection
 - Who has to manage process states? A parent process or kernel?
?

UNIX Process Management APIs



```
int pid = fork();
```

Create a new process that is a clone of its parent.

```
exec*("program" [, argvp, envp]);
```

Overlay the calling process virtual memory with a new program, and transfer control to it.

```
exit(status);
```

Exit with status, destroying the process.

```
int pid = wait*(&status);
```

Wait for exit (or other status change) of a child.

Question: What does this code print?

```
int child_pid = fork();  
if (child_pid == 0) {           // I'm the child process  
    printf("I am process # %d\n", getpid()); — ①  
    return 0;  
} else {                       // I'm the parent process  
    printf("I am parent of process # %d\n", child_pid);  
    return 0; — ②  
}
```

① → ②
② → ① } → don't know, depends on scheduler

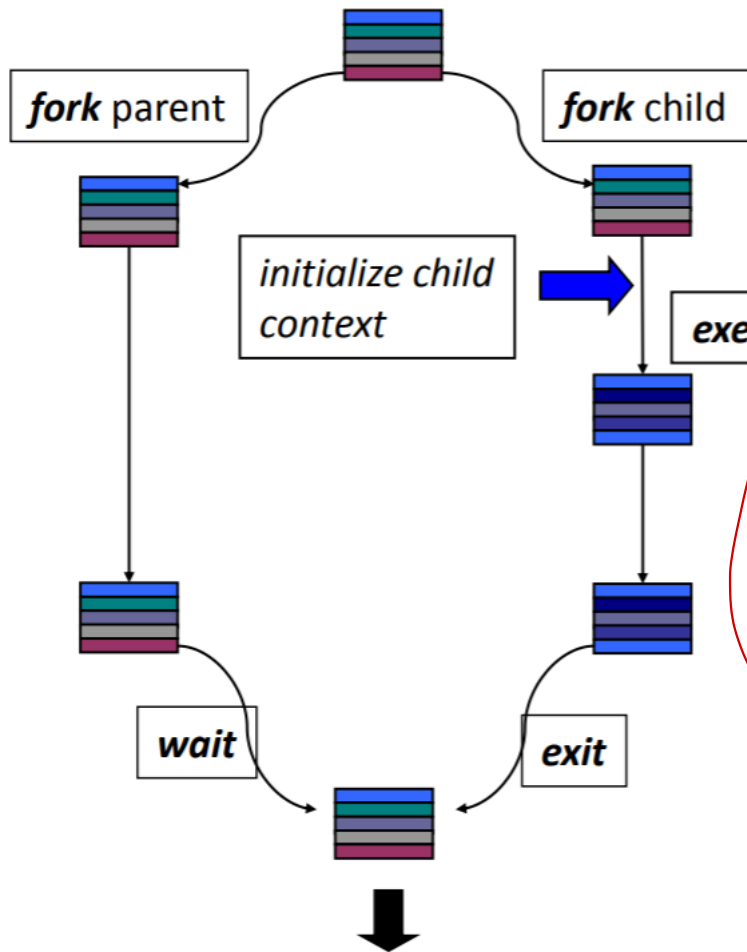
Questions

- Can UNIX fork() return an error? Why?
ex. Too many processes, No memory to allocate
- Can UNIX exec() return an error? Why?
ex. No binary file to execute
- Can UNIX wait() ever return immediately? Why?

Child is already terminated

*Error → Reported to user applications
"errno"*

UNIX Process Management APIs



\$ cc -c sourcefile1.c

parent
Bash (\$) → Bash (child)
→ Run compiler (cc) ← exec
→ [compile] ↓
→ exit
→ return to Bash (\$) ← parent wait


Does execution path above match APIs?

Question: Why fork?

What if just run compile in the same process as bash (\$)?

bash ↔ compile
both don't believe each other
violate protection (about person)
between bash & child

Implementing UNIX fork



Steps to implement UNIX fork

- Create and initialize the process control block (PCB) in the kernel
- Create a new address space to protect between parent & child
- Initialize the address space with a copy of the entire contents of the address space of the parent
- Inherit the execution context of the parent (e.g., any open files)
- Inform the scheduler that the new process is ready to run

Copy on write

Not exactly copy, optimization

Implementing UNIX exec

- Steps to implement UNIX exec
 - Load the program into the current address space
 - Copy arguments into memory in the address space
 - Initialize the hardware context to start execution at ``start''
Application executable file
→ ELF binary format

Implementing a Shell

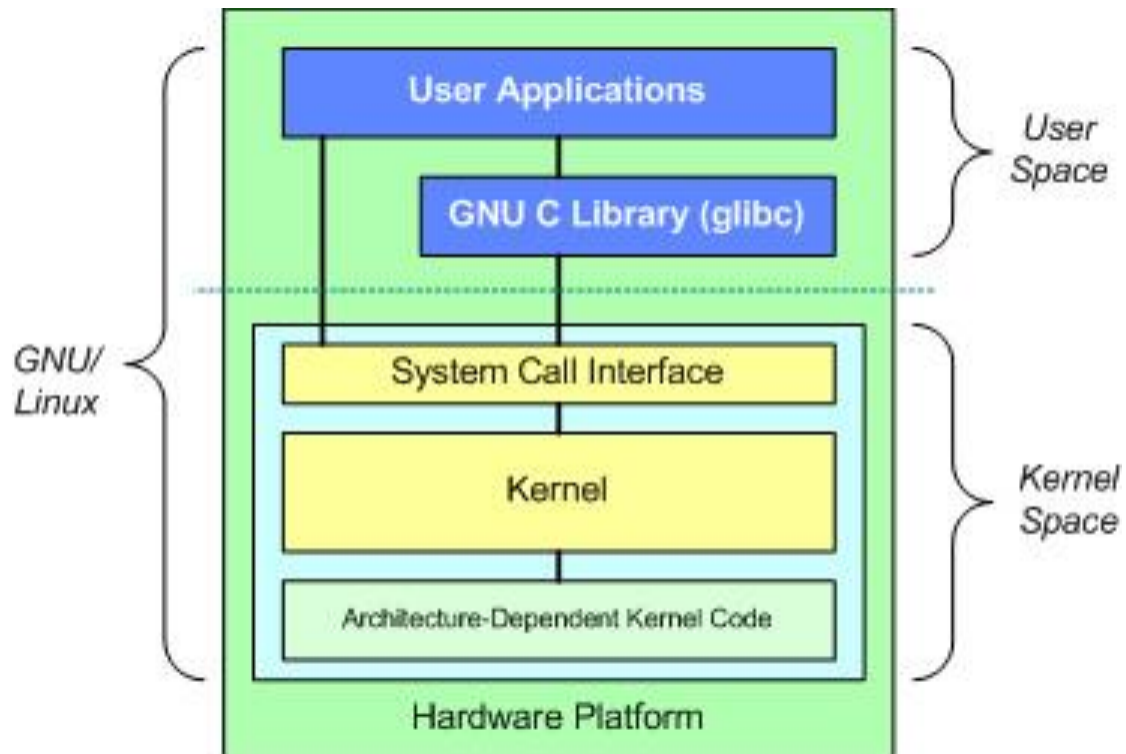
```
char *prog, **args;
int child_pid;

// Read and parse the input a line at a time
while (readAndParseCmdLine(&prog, &args)) {
    child_pid = fork();    // create a child process
    if (child_pid == 0) {
        exec(prog, args);    // I'm the child process. Run program
        // NOT REACHED
    } else {
        wait(child_pid);    // I'm the parent, wait for child
        return 0;
    }
}
```

What to consider designing APIs?

- What functionality should be in kernel?
→ Protection
 - **What functionality should be in user library?**
→ Illusion, Abstraction
Standard libc
- How should OS be organized?

Standard C library



Why is the standard library required?

Which goes which registers? → Protocol

: Application Binary Interface (ABI)

A system call allows a program to request a service—for example, open a file or create a new process—from the kernel. At the assembler level, making a system call requires the caller to assign the unique system call number and the argument values to particular registers, and then execute a special instruction (e.g., SYSENTER on modern x86 architectures) that switches the processor to kernel mode to execute the system-call handling code. Upon return, the kernel places the system call's result status into a particular register and executes a special instruction (e.g., SYSEXIT on x86) that returns the processor to user mode. The usual convention for the result status is that a non-negative value means success, while a negative value means failure. A negative result status is the negated error number (`errno`) that indicates the cause of the failure.

All of the details of making a system call are normally hidden from the user by the C library, which provides a corresponding wrapper function and header file definitions for most system calls. The wrapper function accepts the system call arguments as function

writes d, buf, ... convenient

libc → ABI → kernel

register, ..., (unconvenient)

<https://lwn.net/Articles/534682/>

What to consider designing APIs?

- What functionality should be in kernel?
- What functionality should be in user library?

→ How should OS be organized?

option 1: All functionalities in a **single kernel**


→ Monolithic

option 2: Functionalities are **partitioned to kernel and several services**


Micro kernel

Monolithic Kernel

- All OS functionality is included in the kernel
 - ✱ – Applications do not have control over resources
- Advantages
 - well-understood, **good performance**
- ✱ • Disadvantages
 - **No protection btw kernel components**
 - Not easily extended/modified
- Example:
 - Windows, BSD, Linux




Operating System Structure



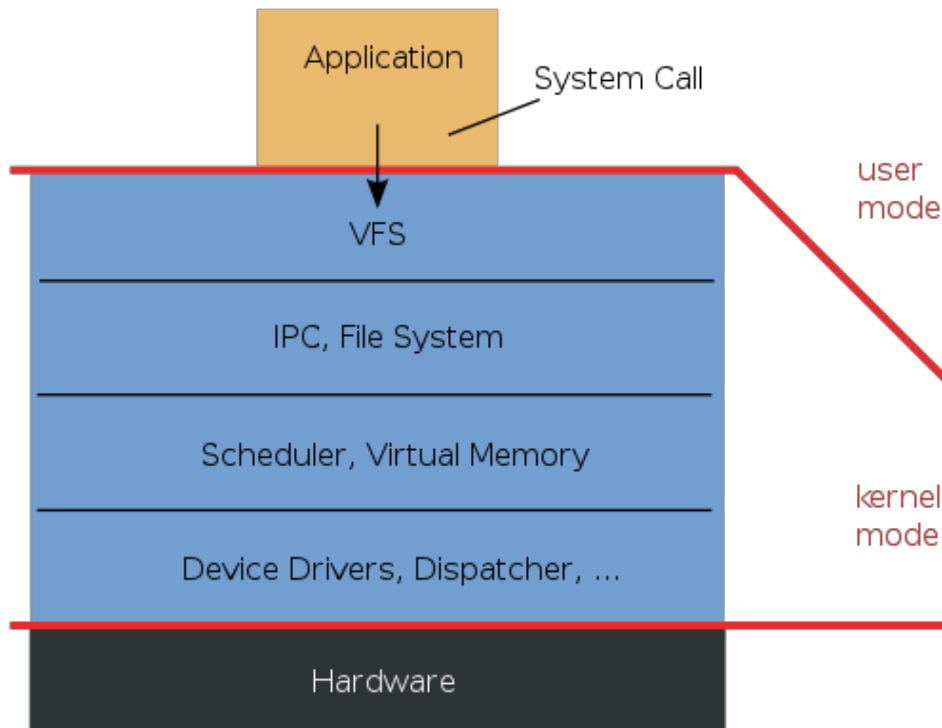
→ Slow: System call requires context switch

-

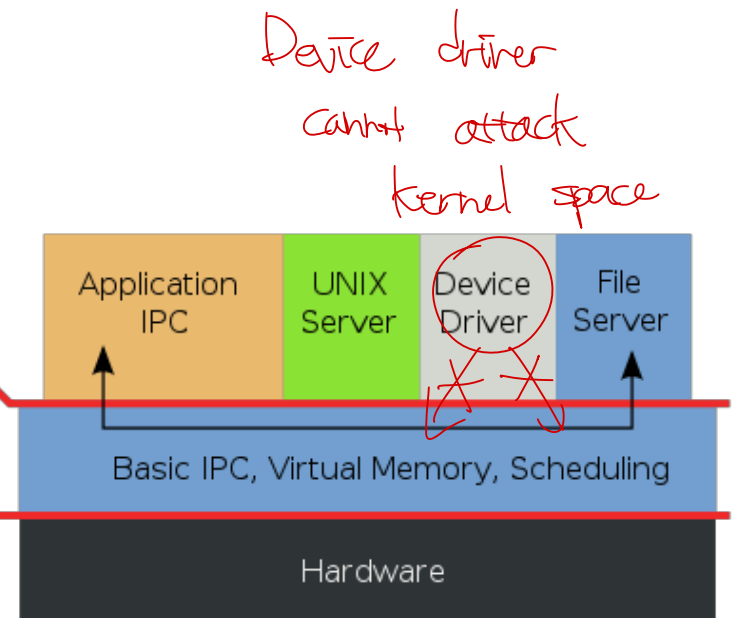


Monolithic Kernels VS Microkernels

Monolithic Kernel based Operating System



Microkernel based Operating System



Microkernel - Advantages

- **Modularity** → Modules in user-level
- **Flexibility and extensibility**
 - Easier to replace modules – fewer dependencies
 - Different servers can implement the same service in different ways
- ★ • **Safety (protection)**
 - Each server is protected by the OS from other servers
 - Servers are isolated; errors in one don't affect others
- **Correctness**
 - Easier to verify a small kernel

Microkernels- Disadvantages

- Performance issue
 - Slow – due to “cross-domain” information transfers?
 - Server-to-OS, OS-to-server IPC is thought to be a major source of inefficiency
 - Much faster to communicate between two modules that are both in OS

Inter-Process Communication

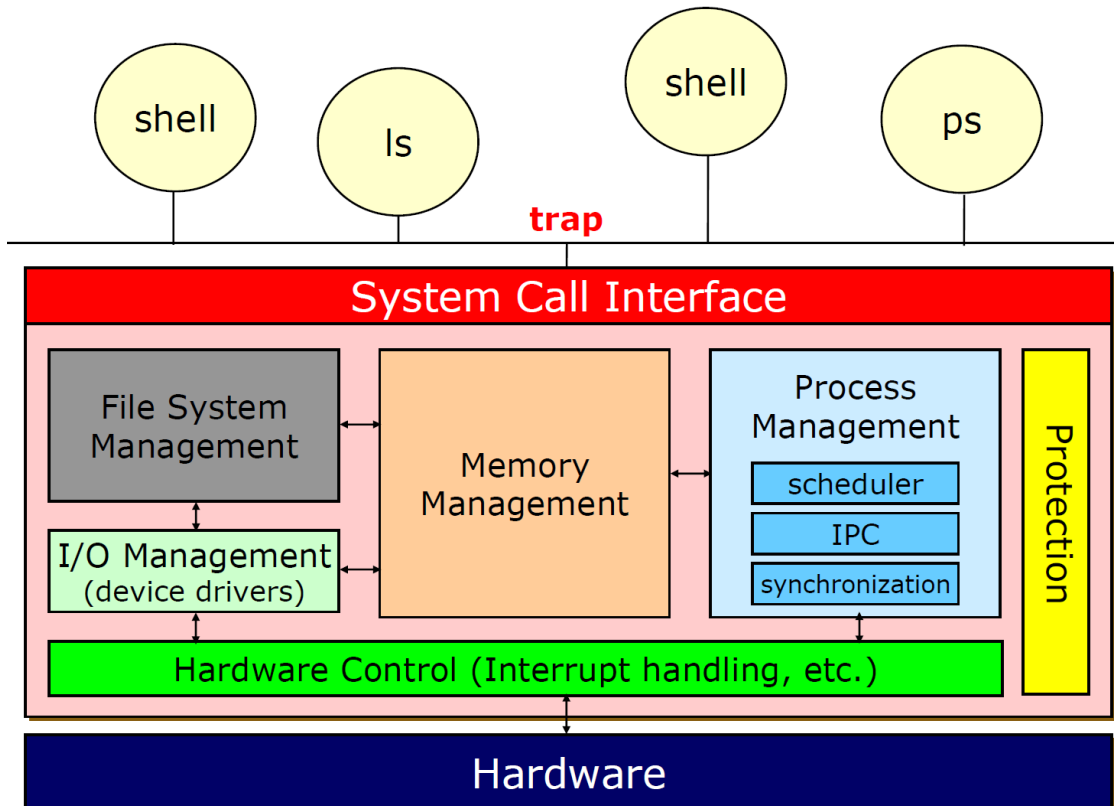
Monolithic OS

Hierarchy of Communication Mechanisms

- Fastest to Slowest
 - Function calls within same process
 - System calls (mode switch) *transition*
 - Context switch (process switch)
 - *↓* **IPC** between processes (message passing, on the same or different machines)

Much slower..

Big picture: where are we?



We've discussed:

- Protection
- Hardware control
- Scheduling
- APIs for process management
- APIs for IO