

CS330: Concurrency and thread

Instructor: Youngjin Kwon

Design: Motivation

- Operating systems (and application programs) often need to be able to handle multiple things happening at the same time
 - Process execution, interrupts, background tasks, system maintenance

Bottleneck




- A phenomenon where the performance of an entire system is limited by one or more components/resources
- System designers will try to avoid bottlenecks
 - try to locate and tune existing bottlenecks



The I/O Bottleneck


- In old days, memory and storage were faster than the processor
 - they were waiting for CPU to finish computation to feed data



The I/O Bottleneck

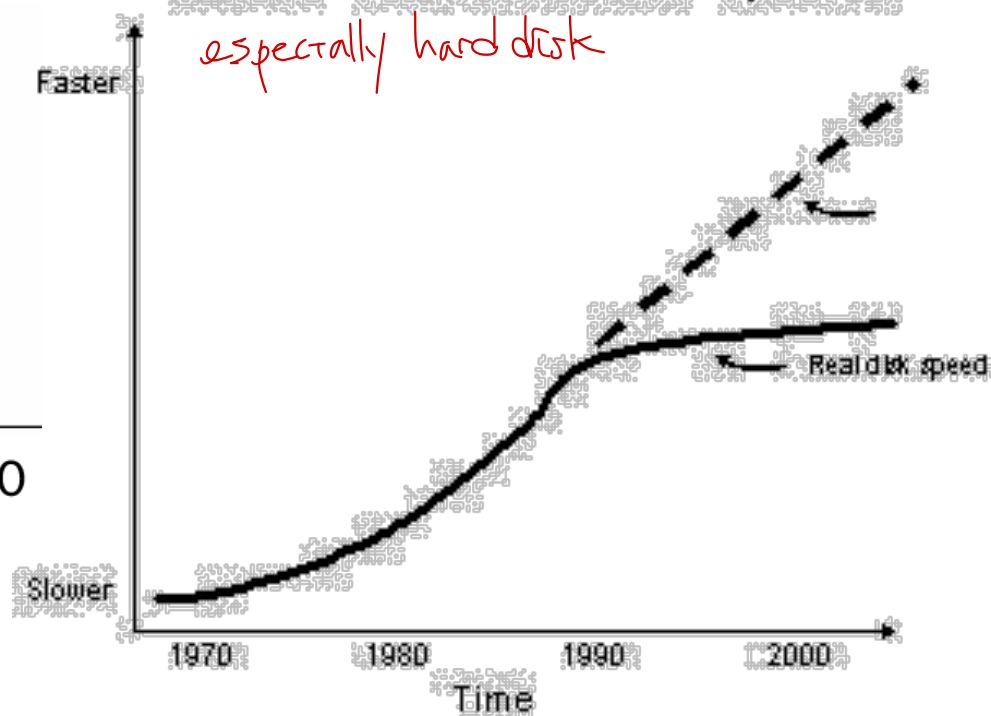
- Over the years, processor speed was improved more than that of memory and storage

Processor Speed




Moore's Law for Disk Speed:

especially hard disk




The I/O Bottleneck

- In old days, memory and storage were faster than the processor
 - they were waiting for CPU to finish computation to feed data



- Over the years, processor speed was improved more than that of memory and storage
- Nowadays, **the processor is waiting** for memory and storage drive to feed data




How to Design OS ?

- How to design OS to overcome the I/O bottleneck?
 - Multi-programming
 - Bigger cache & better cache management algorithms

Multi process

Multi ?



Uniprogramming vs. Multiprogramming

- Uniprogramming: *one program (application) at a time*
 - MS/DOS, early Macintosh, batch processing
 - Easier for operating system builder
 - Get rid of concurrency (only one program accessing resources!)
 - Does this make sense for personal computers?
- Multiprogramming: *more than one programs at a time*
 - Multics, UNIX/Linux, OS/2, Windows NT – 7, Mac OS X

Solution Design


Goals:

Addressing the I/O bottleneck



Solution:

- Multiprogramming: Run multiple applications concurrently



How to design?

Design: Need a new abstraction!

a schedulable task of execution stream

- A **thread** is a **single execution sequence** that represents a **separately schedulable task**
 - Single execution sequence: familiar programming model
 - Separately schedulable: OS can run or suspend a thread at any time
- **Protection is an orthogonal concept**
 - Can have one or many threads per protection domain

- Process

- **Protection unit**

- Abstraction of (*Machine*)

CPU + memory + I/O



- Thread

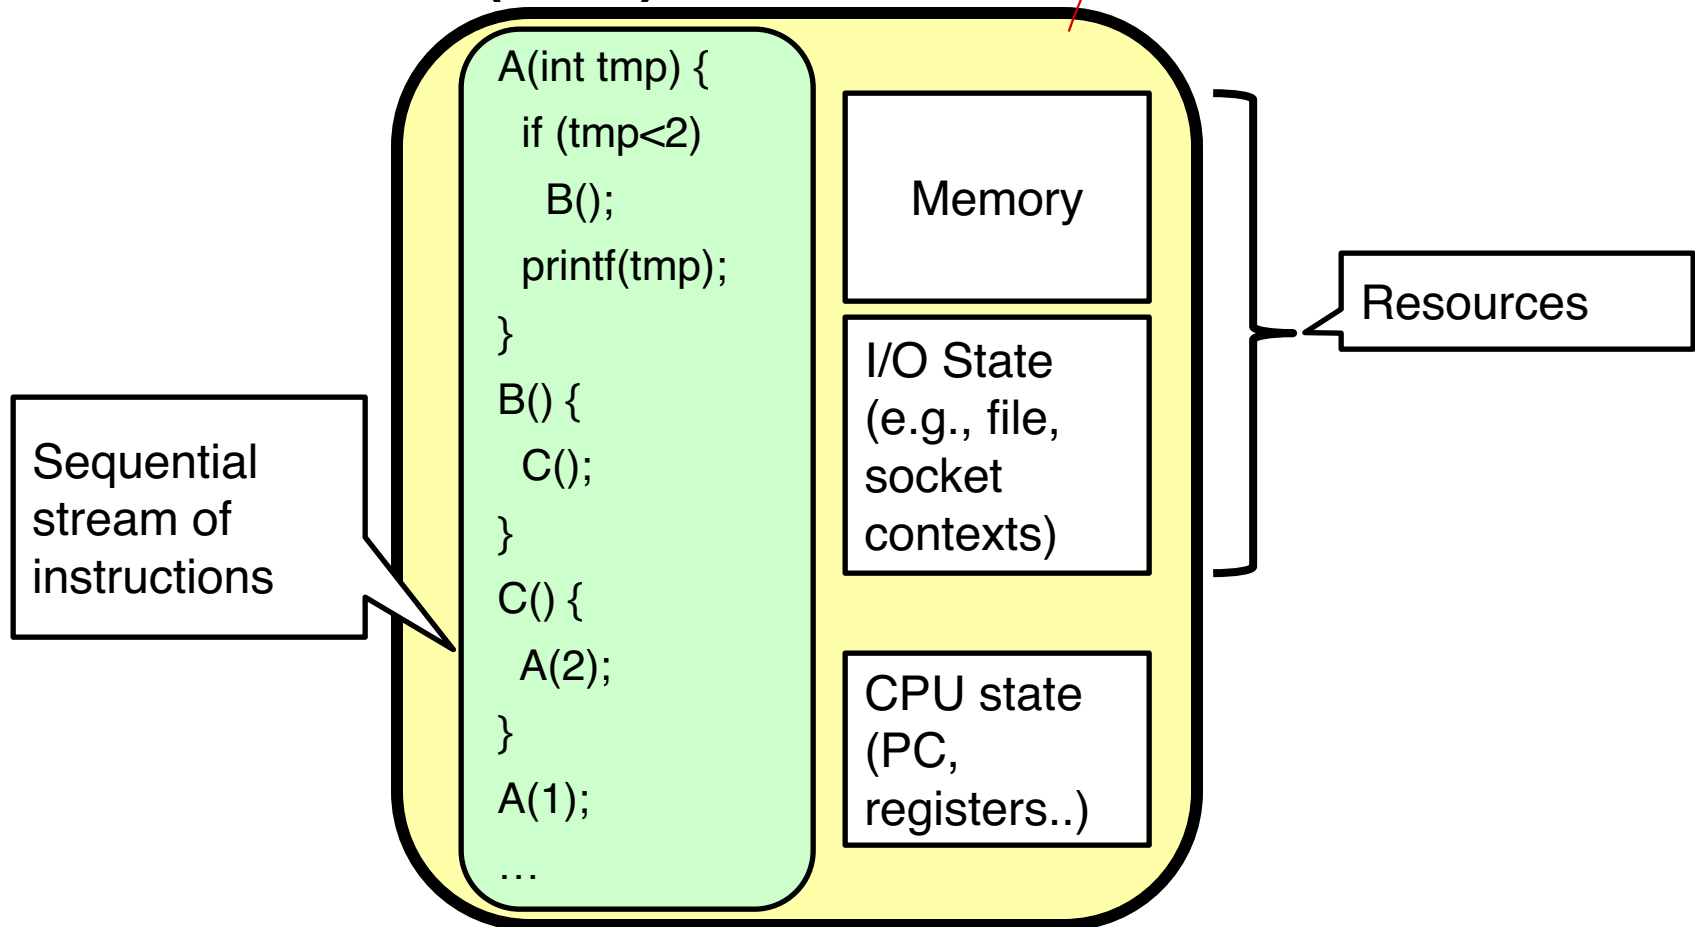
- **Execution unit**

- Abstraction of (*CPU*)

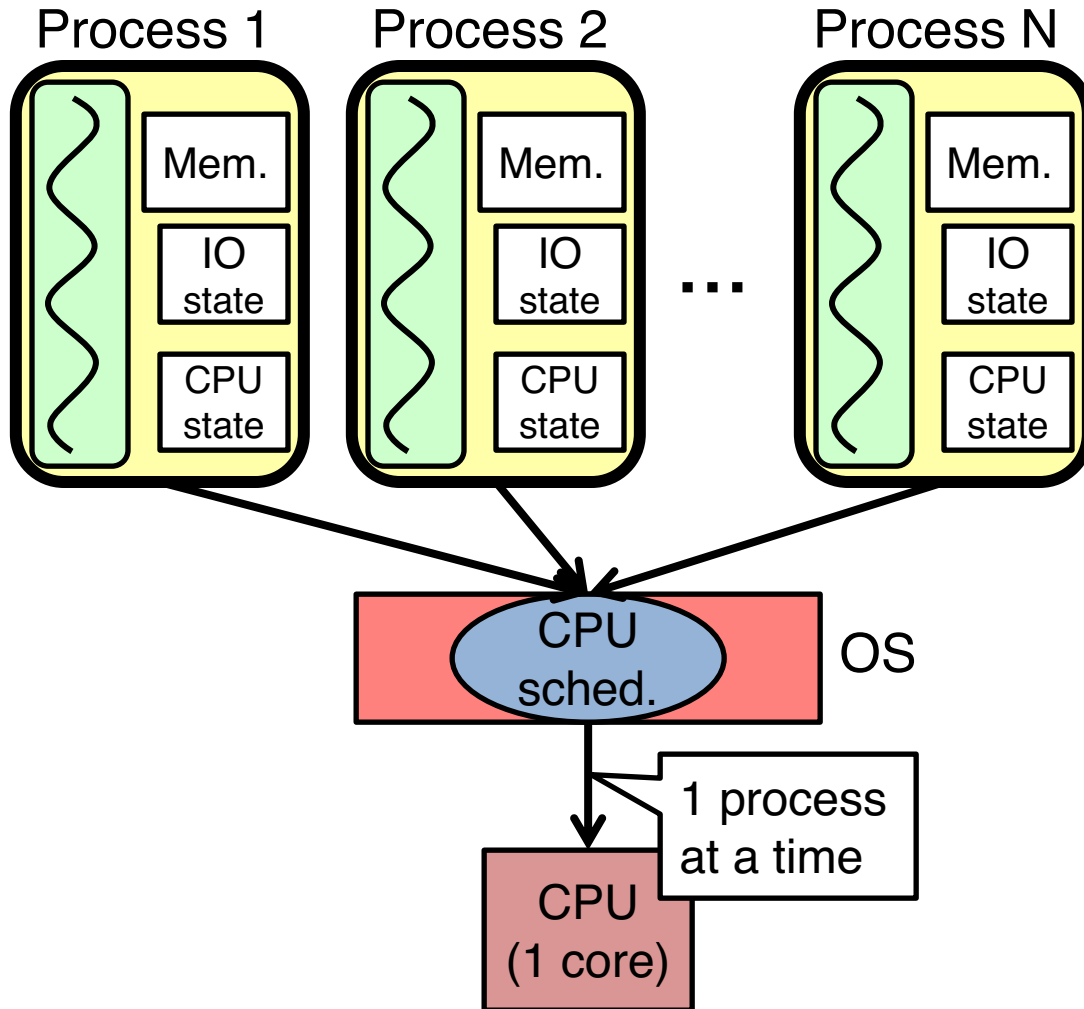
Process

(Unix) Process

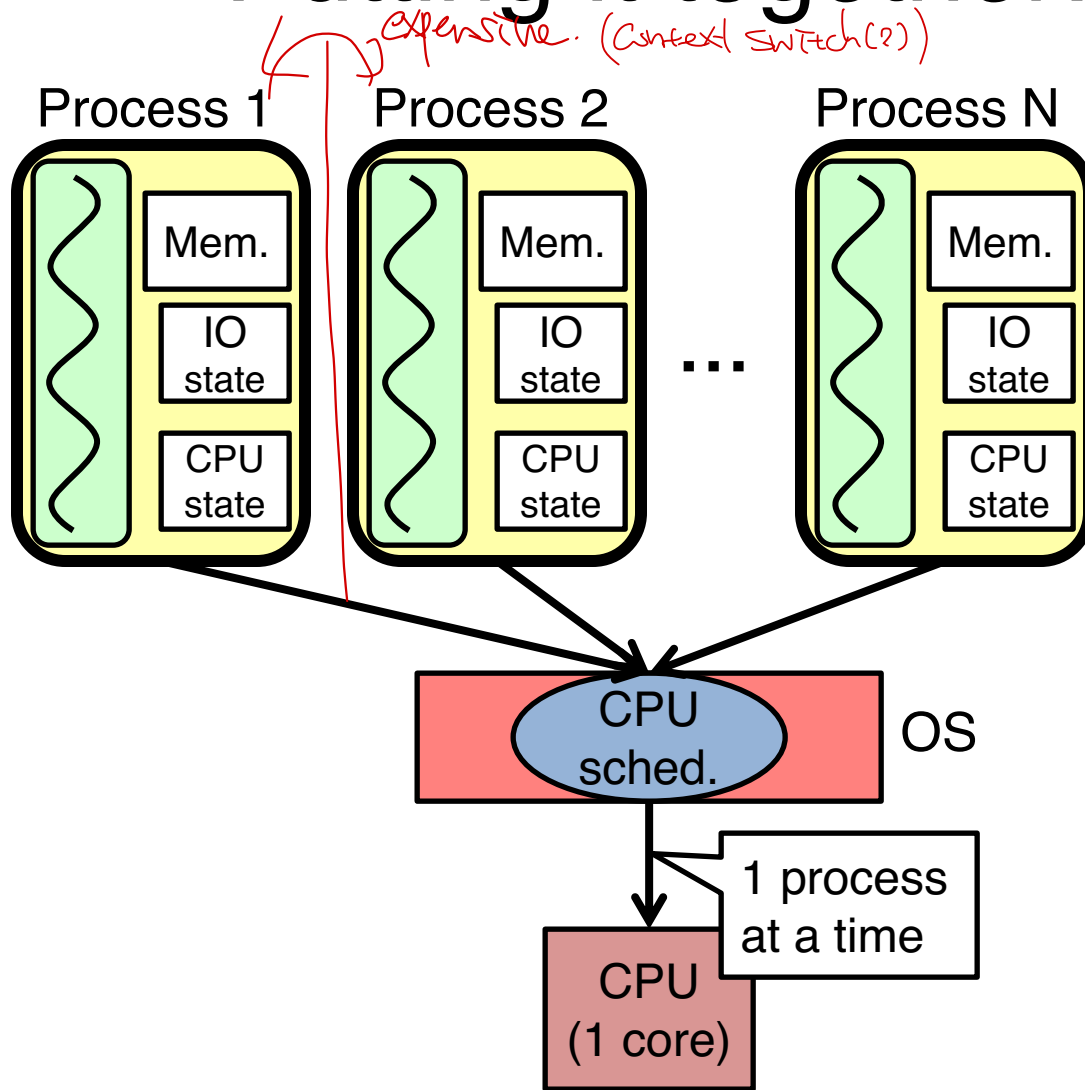
PCB (Process Control Block)



Processes

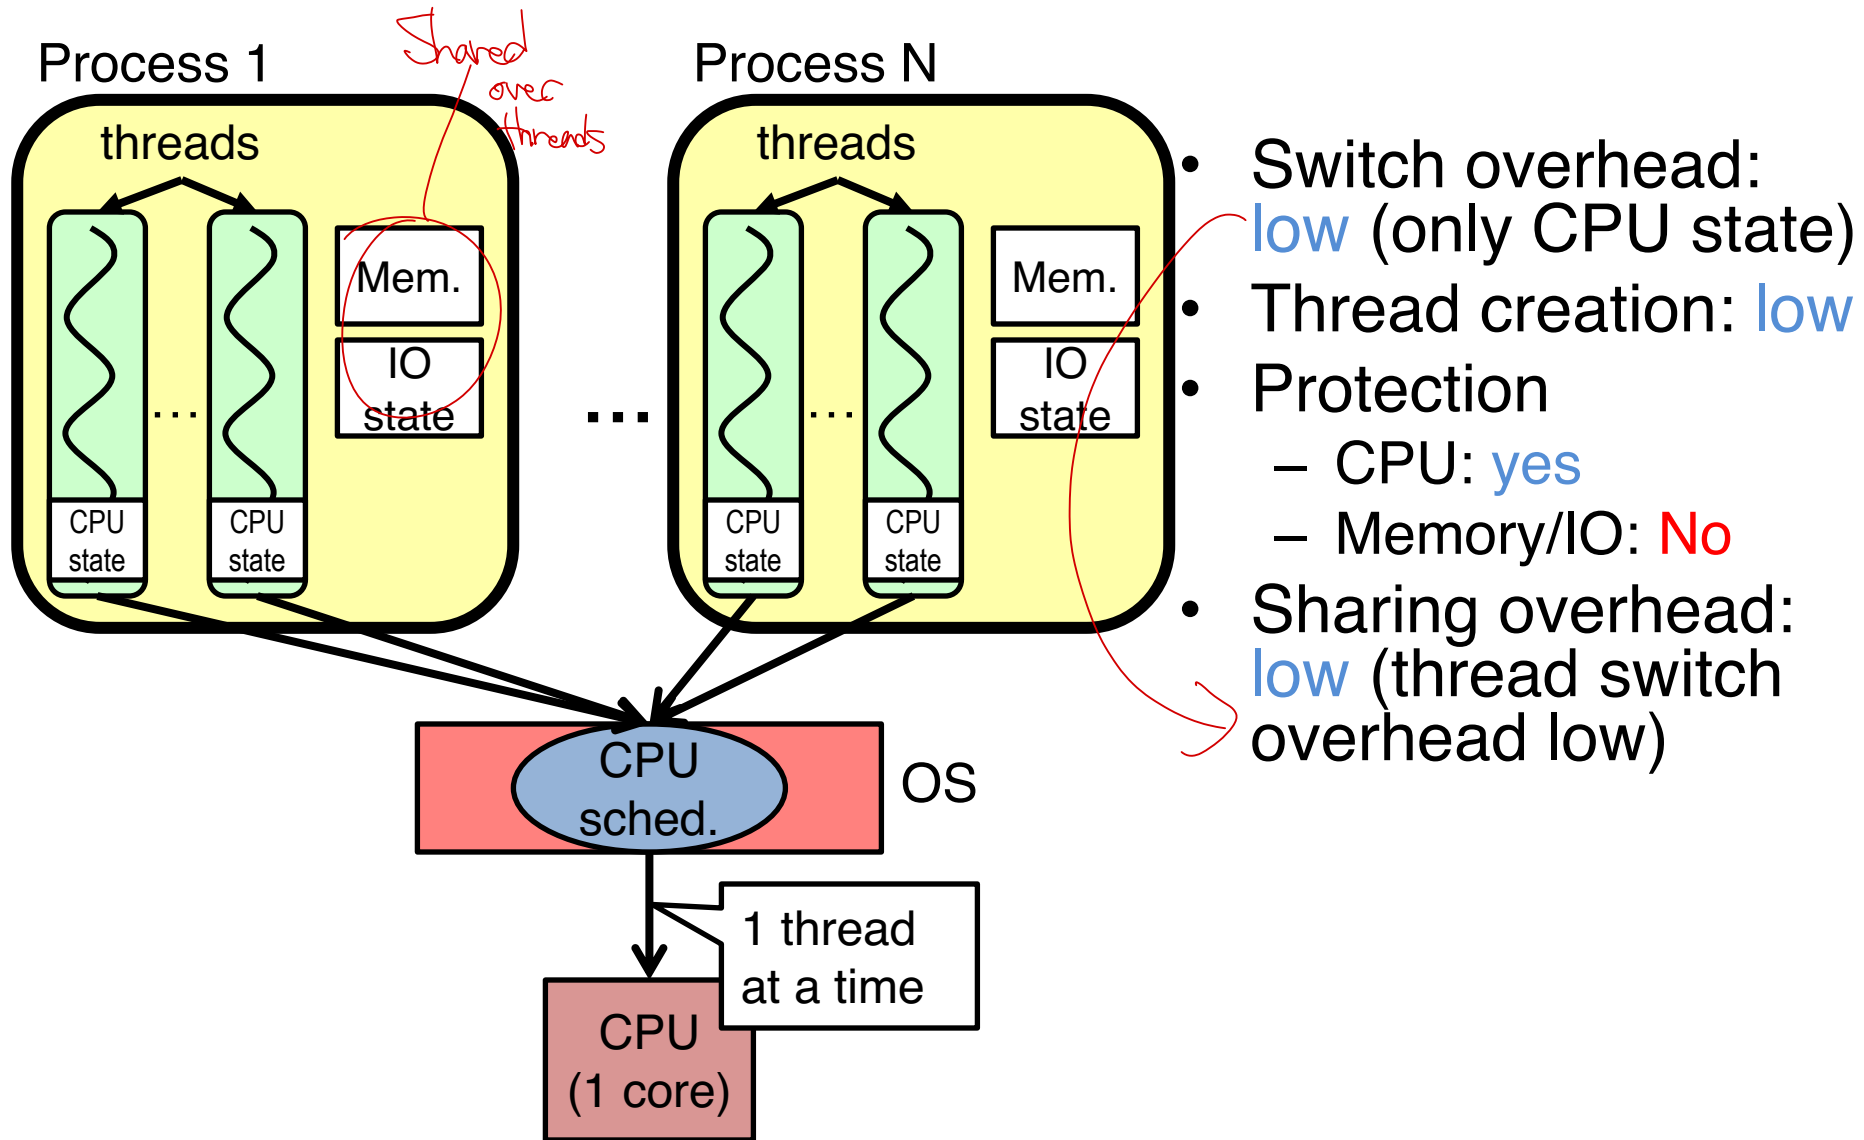


Putting it together: Processes



- **Switch overhead: high**
 - CPU state: **low**
 - Memory/IO state: **high**
 - Process creation: **high**
 - Protection
 - CPU: **yes**
 - Memory/IO: **yes**
 - Sharing overhead: **high** (involves at least a context switch)
- Handwritten notes:* "few registers" (with an arrow pointing to the "CPU state" box of Process 1), "Page table alloc" (with an arrow pointing to the "Change Page Table" note).

Putting it together: Threads



Review: Execution Stack Example

addrX:	A(int tmp) {
.	if (tmp<2)
.	B();
addrY:	printf(tmp);
.	}
.	B() {
.	C();
addrU:	}
.	C() {
.	A(2);
addrV:	}
.	A(1);
.	
addrZ:	exit;

- Stack holds function arguments, return address → *Compiler*
- Permits recursive execution
- Crucial to modern languages

Review: Execution Stack Example

addrX:	A(int tmp) {
.	if (tmp<2)
.	B();
addrY:	printf(tmp);
.	}
.	B() {
.	C();
addrU:	}
.	C() {
.	A(2);
addrV:	}
.	A(1);
addrZ:	exit;

PC

Stack
Pointer

A: tmp=1
ret=addrZ

Stack Growth

- Stack holds function arguments, return address
- Permits recursive execution
- Crucial to modern languages

Review: Execution Stack Example

addrX:	A(int tmp) {
.	if (tmp<2)
.	B();
addrY:	printf(tmp);
.	}
.	B() {
.	C();
addrU:	}
.	C() {
.	A(2);
addrV:	}
.	A(1);
addrZ:	exit;

PC

Stack
Pointer


A: tmp=1
ret=addrZ

Stack Growth

- Stack holds function arguments, return address
- Permits recursive execution
- Crucial to modern languages

Review: Execution Stack Example

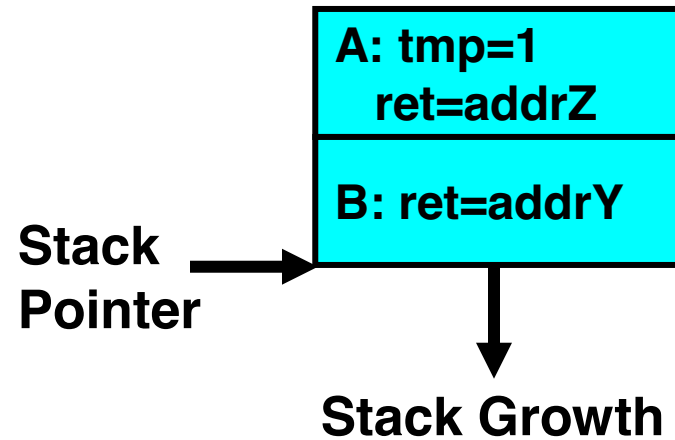
addrX:	A(int tmp) {
.	if (tmp<2)
.	
.	B();
addrY:	printf(tmp);
.	}
.	
.	B() {
.	C();
addrU:	}
.	
.	C() {
.	A(2);
addrV:	}
.	
.	A(1);
addrZ:	exit;



- Stack holds function arguments, return address
- Permits recursive execution
- Crucial to modern languages

Review: Execution Stack Example


addrX:	A(int tmp) {
.	if (tmp<2)
.	B();
addrY:	printf(tmp);
.	}
.	B() {
.	C();
addrU:	}
.	C() {
.	A(2);
addrV:	}
.	A(1);
.	
addrZ:	exit;



- Stack holds function arguments, return address
- Permits recursive execution
- Crucial to modern languages

Review: Execution Stack Example

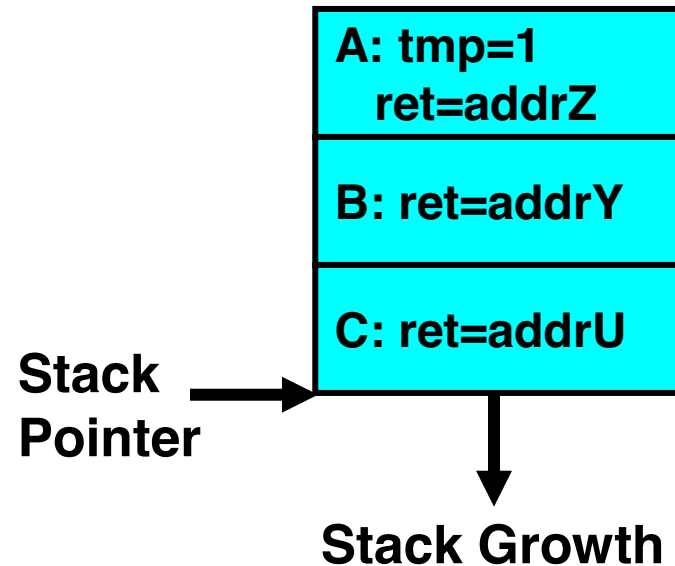
addrX:	A(int tmp) {
.	if (tmp<2)
.	B();
addrY:	printf(tmp);
.	}
.	B() {
.	C();
addrU:	}
.	C() {
.	A(2);
addrV:	}
.	A(1);
addrZ:	exit;



- Stack holds function arguments, return address
- Permits recursive execution
- Crucial to modern languages

Review: Execution Stack Example

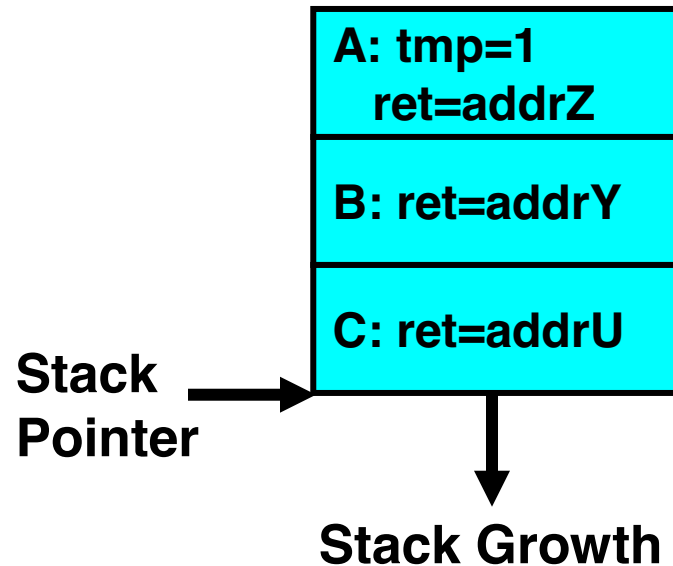
addrX:	A(int tmp) {
.	if (tmp<2)
.	B();
addrY:	printf(tmp);
.	}
.	B() {
.	C();
addrU:	}
.	C() {
.	A(2);
addrV:	}
.	A(1);
addrZ:	exit;



- Stack holds function arguments, return address
- Permits recursive execution
- Crucial to modern languages

Review: Execution Stack Example

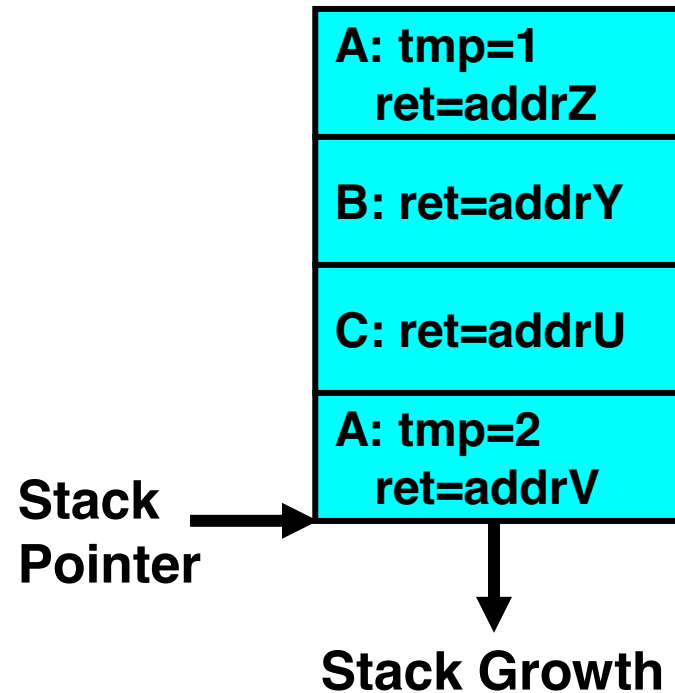
addrX:	A(int tmp) {
.	if (tmp<2)
.	B();
addrY:	printf(tmp);
.	}
.	B() {
.	C();
addrU:	}
.	C() {
.	A(2);
addrV:	}
.	A(1);
addrZ:	exit;



- Stack holds function arguments, return address
- Permits recursive execution
- Crucial to modern languages

Review: Execution Stack Example


addrX:	A(int tmp) {
.	if (tmp<2)
.	B();
addrY:	printf(tmp);
.	}
.	B() {
.	C();
addrU:	}
.	C() {
.	A(2);
addrV:	}
.	A(1);
addrZ:	exit;



- Stack holds function arguments, return address
- Permits recursive execution
- Crucial to modern languages

Review: Execution Stack Example

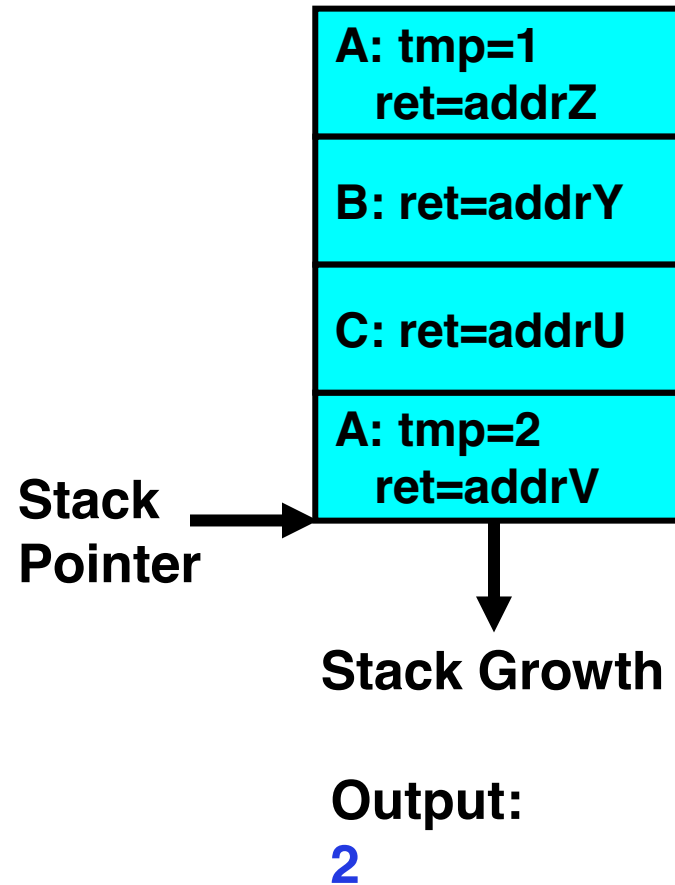
addrX:	A(int tmp) {
.	if (tmp<2)
.	
.	B();
addrY:	printf(tmp);
.	}
.	B() {
.	C();
addrU:	}
.	C() {
.	A(2);
addrV:	}
.	A(1);
.	
addrZ:	exit;



- Stack holds function arguments, return address
- Permits recursive execution
- Crucial to modern languages

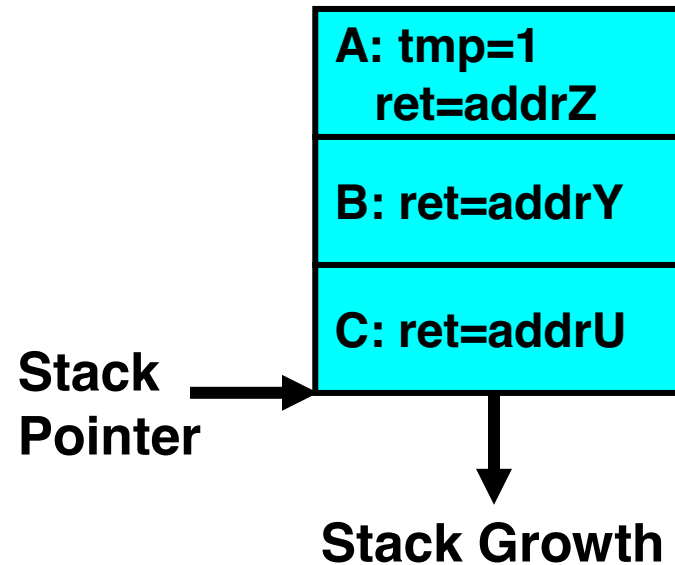
Review: Execution Stack Example

addrX:	A(int tmp) {
.	if (tmp<2)
.	B();
addrY:	printf(tmp);
.	}
.	B() {
.	C();
addrU:	}
.	C() {
.	A(2);
addrV:	}
.	A(1);
addrZ:	exit;



Review: Execution Stack Example

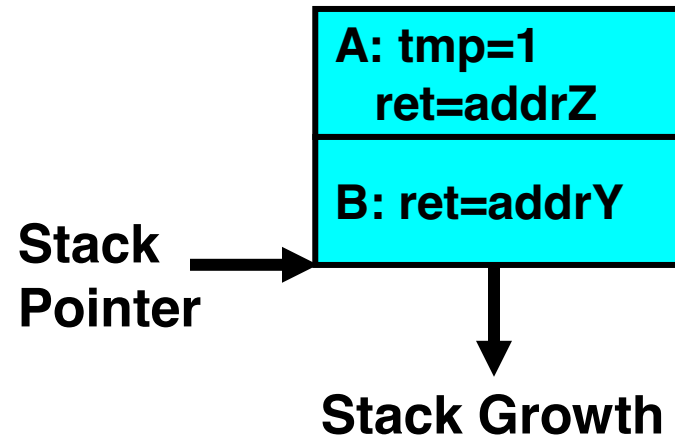
```
addrX:  A(int tmp) {  
        .  
        .    if (tmp<2)  
        .        B();  
addrY:  printf(tmp);  
        .  
        .    }  
        .  
        .    B() {  
        .        C();  
addrU:  }  
        .  
        .    C() {  
        .        A(2);  
addrV:  }  
        .  
        .    A(1);  
addrZ:  exit;
```



Output:
2

Review: Execution Stack Example


addrX:	A(int tmp) {
.	if (tmp<2)
.	B();
addrY:	printf(tmp);
.	}
.	B() {
.	C();
addrU:	}
.	C() {
.	A(2);
addrV:	}
.	A(1);
addrZ:	exit;



Output:
2

Review: Execution Stack Example

addrX:	A(int tmp) {
.	if (tmp<2)
.	B();
addrY:	printf(tmp);
.	}
.	B() {
.	C();
addrU:	}
.	C() {
.	A(2);
addrV:	}
.	A(1);
addrZ:	exit;



Output:

2
1

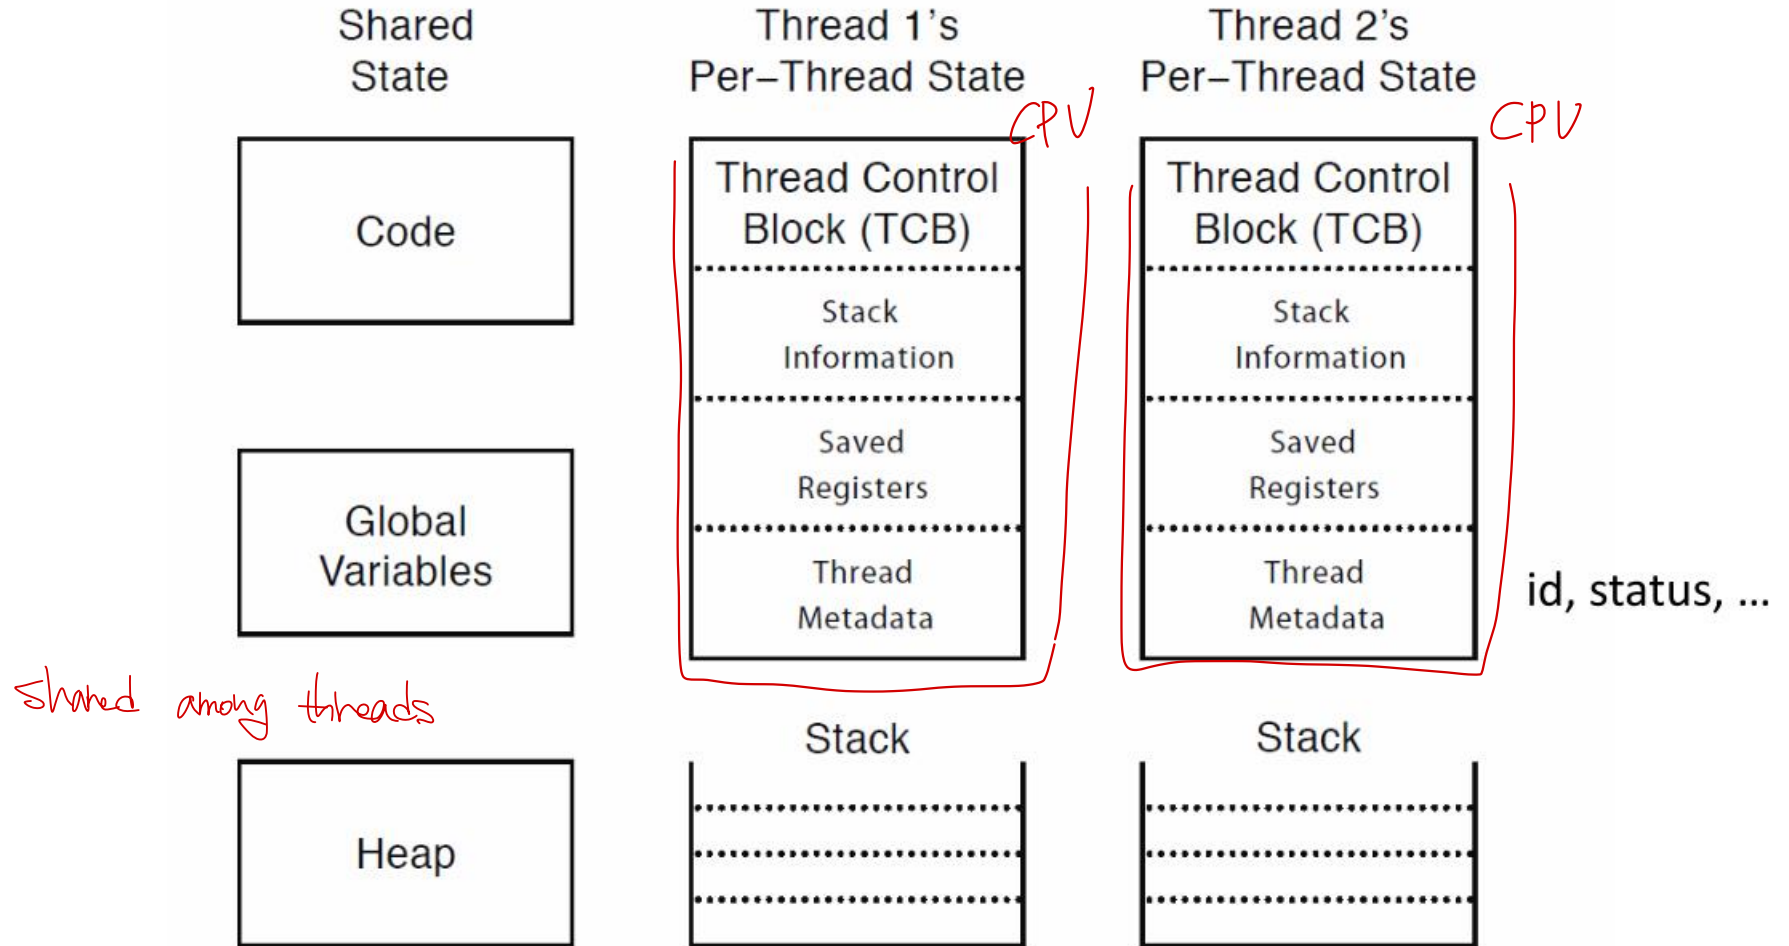
Review: Execution Stack Example

addrX:	A(int tmp) {
.	if (tmp<2)
.	B();
addrY:	printf(tmp);
.	}
.	B() {
.	C();
addrU:	}
.	C() {
.	A(2);
addrV:	}
.	A(1);
addrZ:	exit;

Output:

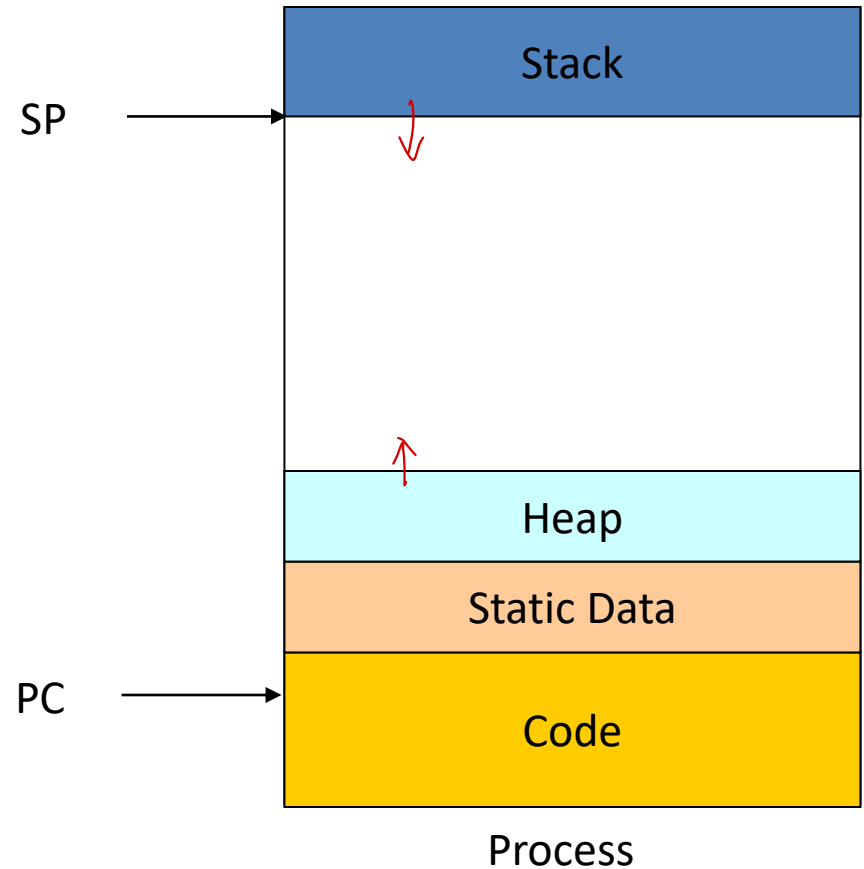
2
1

Thread Data Structures



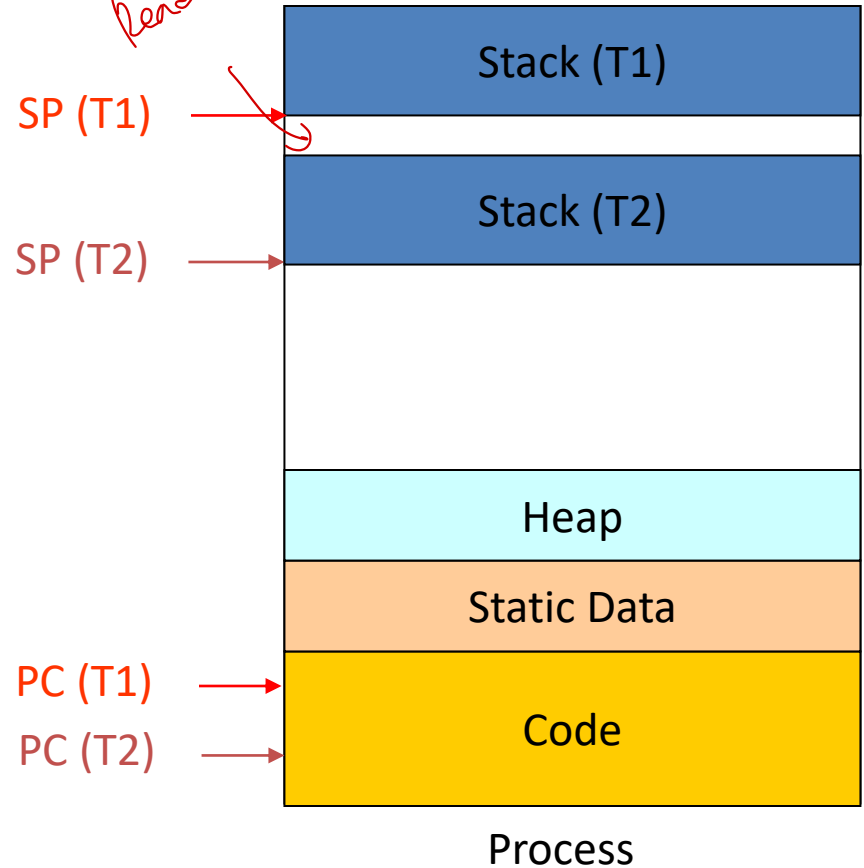
Process vs. Thread

- Execution context
 - Program counter (PC)
 - Stack pointer (SP)
 - Data registers
- Code
- Data
- Stack




Process vs. Thread

- Execution context
 - Program counter (PC)
 - Stack pointer (SP)
 - Data registers
- Stack




Possible Executions

One Execution




Single-core, non-preempt

Another Execution



preempt, time limit, cancel

Another Execution



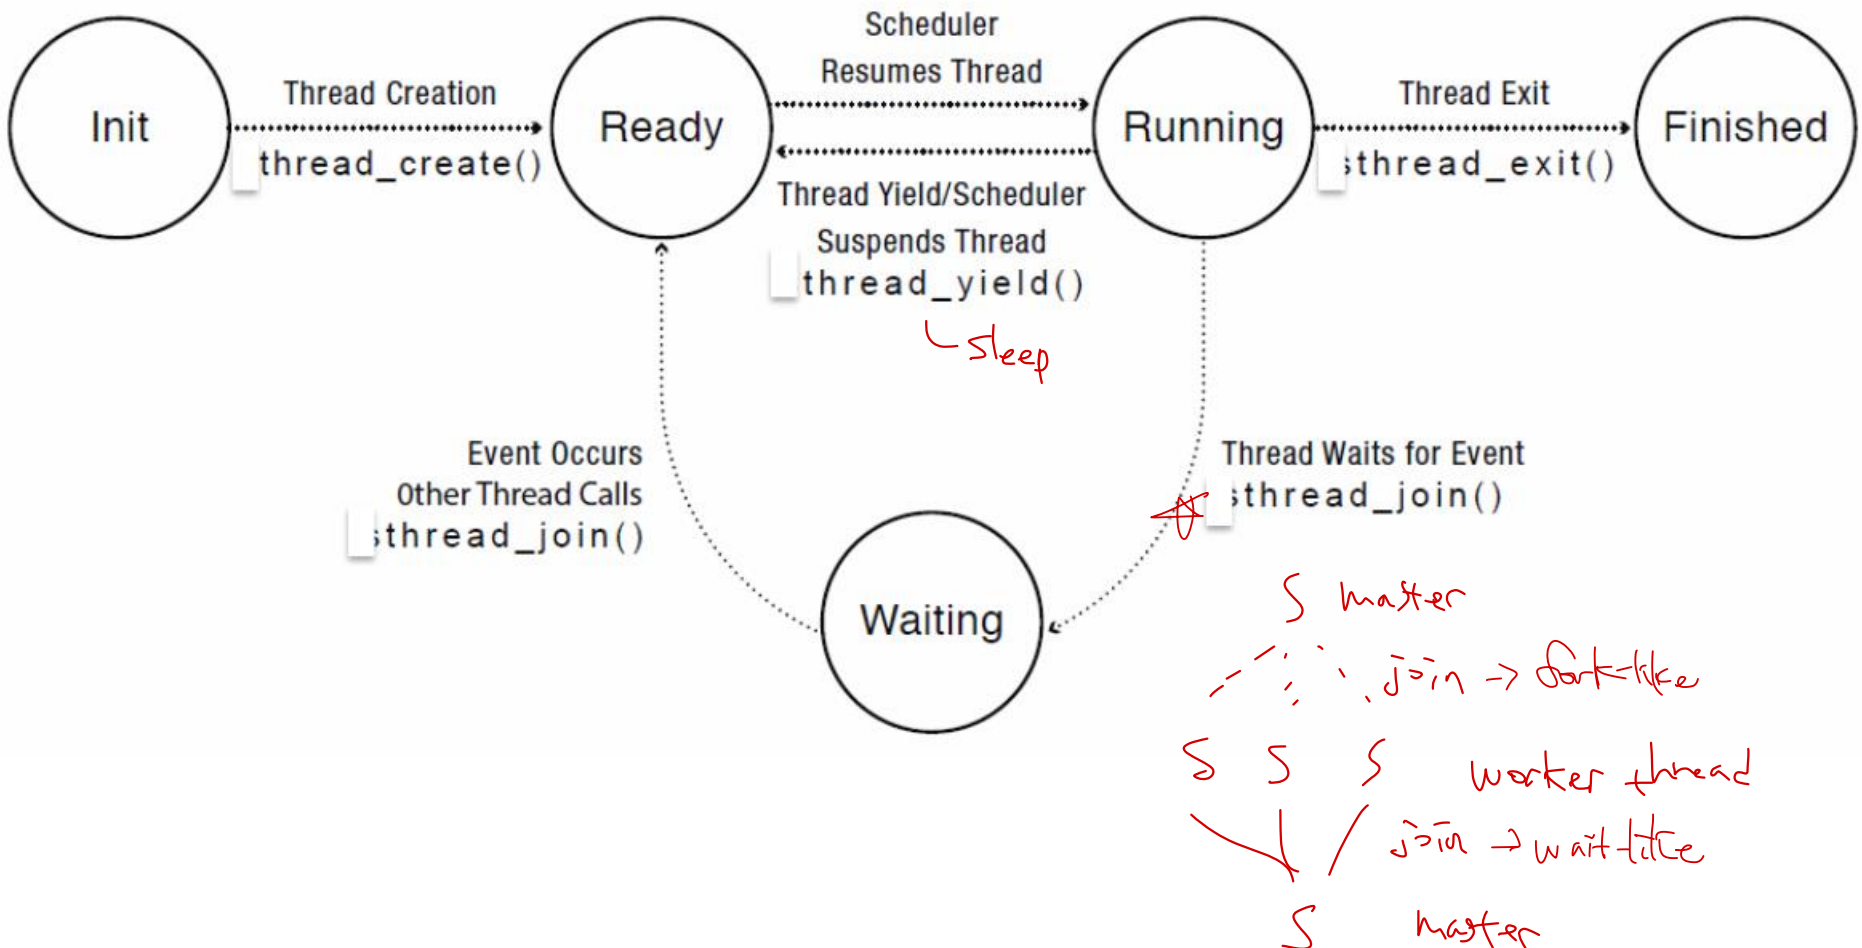
Multicore

Thread performance is highly unpredictable!

Thread programming model

- Cannot assume execution time of other threads (which one is correct?)
 - ✗ – Thread 1 uses sleep (3) to wait for thread 2 to finish its task → never expect timing of other threads
 - ⊖ – Thread 1 waits a signal from thread 2 to wait for thread 2 to finish its task → Semaphore, Conditional variable
- Need **synchronization** when accessing shared data
 - Synchronization provides^{||} **expected execution orders**^{||}

Thread Lifecycle to Design APIs



Thread APIs

- `thread_create(thread, func, args)`
 - Create a new thread to run `func(args)`
- `thread_yield()`
 - Relinquish^{exit} processor voluntarily
- `thread_join(thread)`
 - In parent, wait for forked thread to exit, then return
- `thread_exit`
 - Quit thread and clean up, wake up joiner if any

Example: threadHello

```
#define NTHREADS 10
thread_t threads[NTHREADS];
main() {
    for (i = 0; i < NTHREADS; i++) thread_create(&threads[i], &go, i);
    for (i = 0; i < NTHREADS; i++) {
        exitValue = thread_join(threads[i]);
        printf("Thread %d returned with %ld\n", i, exitValue);
    }
    printf("Main thread done.\n");
}

void go (int n) {
    printf("Hello from thread %d\n", n);
    thread_exit(100 + n);
    // REACHED?
}
```

Create 10 threads

→ All threads are done ✓

but exec order not guaranteed

Design: What is the next?

- Now, you built the thread abstraction
 - How the thread look like
 - Execution model of the thread
 - Define necessary data structures for thread
- Then?

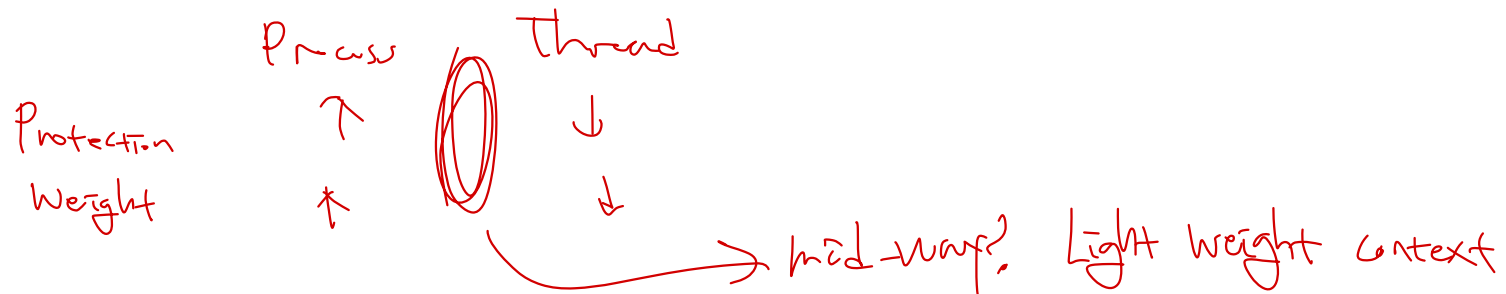
Schedule

Thread scheduling

- Concept of scheduler
 - Map execution unit to processor
 - How to (policies)?
 - FCFS, SJF, RR, priority scheduling, and MLFQ
 - We already covered the scheduling policies

Summary: Process vs. Thread

- A thread is bounded to a single process
- Processes are now containers in which threads execute
- A process can have multiple threads
- Sharing data between threads is cheap: all see the same address space
- Threads become the unit of scheduling



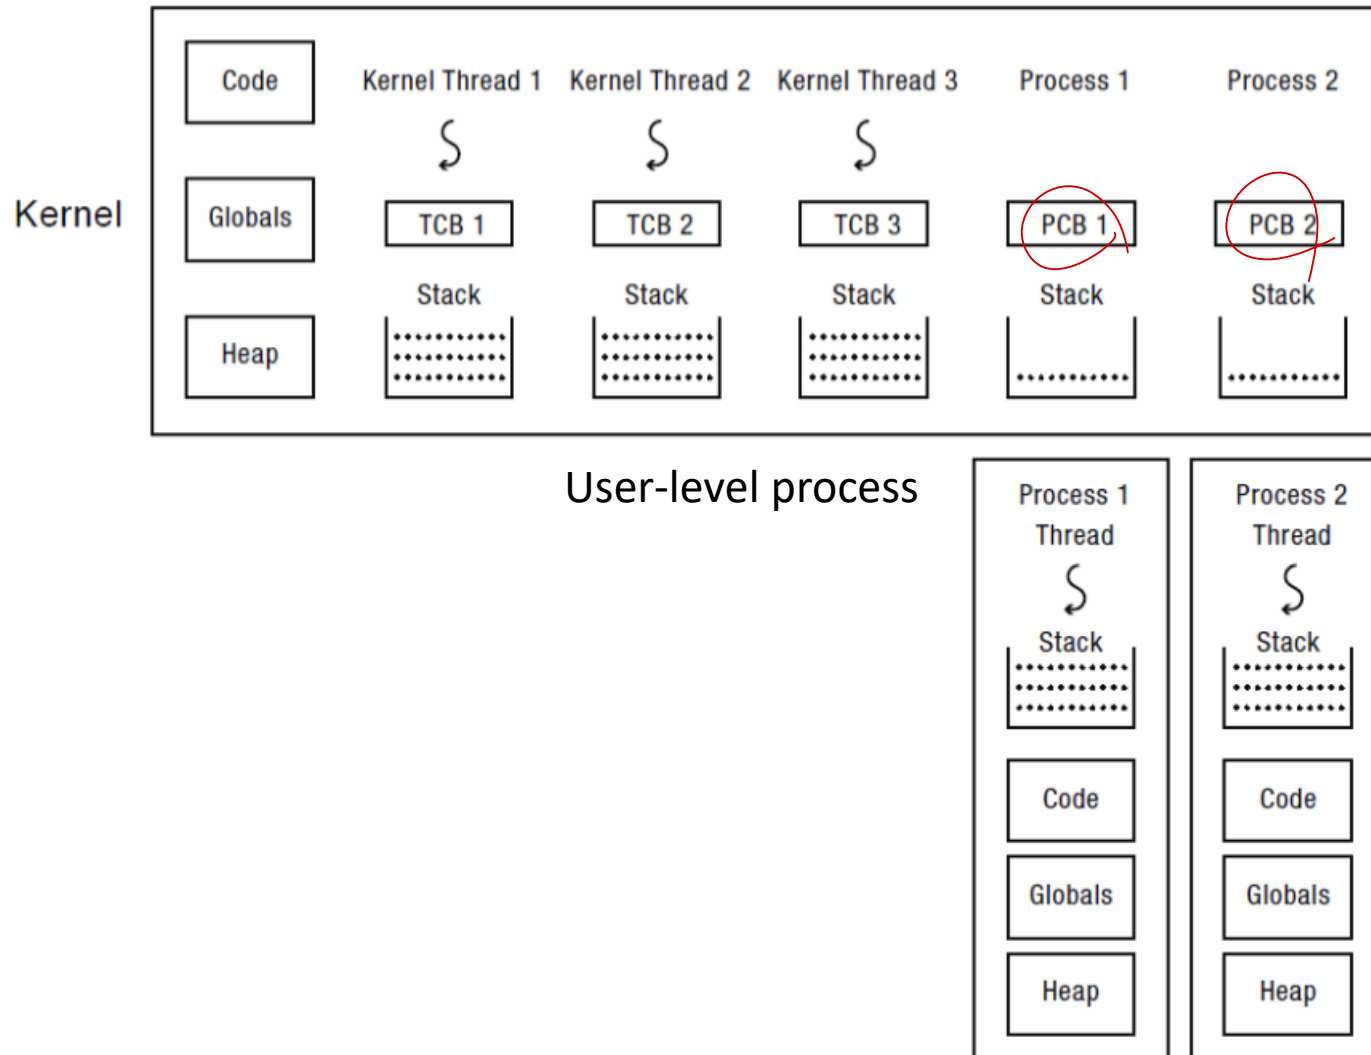
Implementation question

- Who takes care of thread management?
 - OS (kernel threads)
 - System calls for thread creation and management
 - User-level process (user-level threads)
 - a library linked into the program takes care of it

Implementing Threads: Roadmap

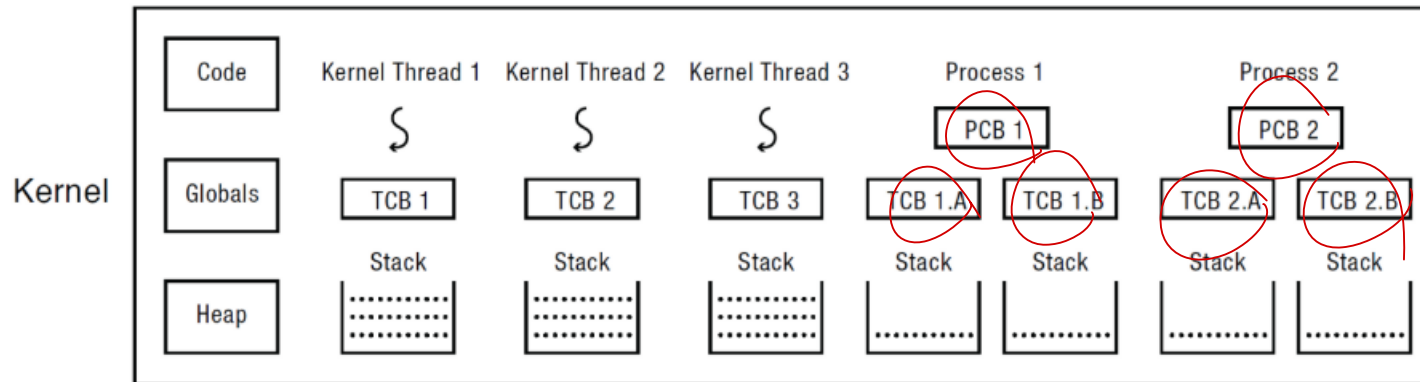
- Kernel threads
 - Thread abstraction only available to kernel
 - To the kernel, a kernel thread and a single threaded user process look quite similar
- Multithreaded processes using kernel threads (Linux, MacOS)
 - Kernel thread operations available via syscall
- User-level threads
 - Thread operations without system calls

Multithreaded OS Kernel single-threaded process

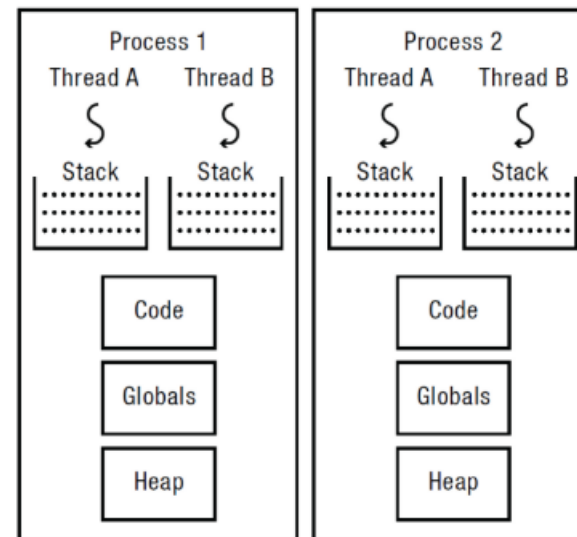


Multithreaded OS Kernel

Multi-threaded process



User-level process



pthread


→ standard thread API

→ implement: kernel thread


Threading model

<https://www.crocus.co.kr/1404>


- Kernel thread




- User-level thread



Implementing threads in Kernel



Implementing threads in user space



Where is thread scheduler?

User thread library

What is performance benefit?

Overhead (No need syscall)

What is a main problem?

Blocked when syscall, I/O

All will

동시에

한 프로그램은 CPU 사용도 못함

User level에서 실행도 안

프로그램 실행

Kernel thread limitations

- Every thread operation must go through kernel
 - create, exit, join, synchronize, or switch for any reason
 - On my laptop: syscall takes 100 cycles, fn call 5 cycles
 - Result: threads 10x-30x slower when implemented in kernel
- Thread context switch overhead↑ (TCB, PCB control)
- One-size fits all thread implementation
 - Kernel threads must please all people
 - Maybe pay for fancy features (priority, etc.) you don't need
- General heavy-weight memory requirements
 - e.g., requires a fixed-size stack within kernel
 - other data structures designed for heavier-weight processes

TCB, PCB 모두 OS가 관리

User-level thread limitations

- Can't take advantage of multiple CPUs or cores
- User-level threads are invisible to the OS
 - They are not well integrated with the OS
- As a result, the OS can make poor decisions
 - Scheduling a process with idle threads
 - A blocking system call blocks all threads
 - Can replace read to handle network connections, but usually OSes don't let you do this for disk
 - Unscheduling a process with a thread holding a lock
- How to solve this?
 - Communication between the kernel and the user-level thread manager (Windows 8) [Scheduler Activation]


Summary

- Kernel-level threads
 - Integrated with OS (informed scheduling)
 - Slower to create, manipulate, synchronize
- User-level threads
 - Faster to create, manipulate, synchronize
 - Not integrated with OS (uninformed scheduling)
- Understanding their differences is important
 - Correctness, performance

Kernel and User threads

- Or use both kernel and user-level threads
 - Can associate a user-level thread with a kernel-level thread
 - Or, multiplex user-level threads on top of kernel-level threads
- **Java Virtual Machine (JVM)** (also C#, others)
 - Java threads are user-level threads
 - On older Unix, only one “kernel thread” per process
 - Multiplex all Java threads on this one kernel thread
 - On modern OSes
 - Can multiplex Java threads on multiple kernel threads
 - Can have more Java threads than kernel threads

User threads on Kernel threads



- User threads implemented on kernel threads
 - Multiple kernel-level threads per process
- Sometimes called **n:m threading**
 - Have n user threads per m kernel threads (user-level threads are n : 1, kernel threads 1 : 1)