


CS330: Synchronization

Instructor: Youngjin Kwon

Threaded Web Server




- Multi-threaded version:

```
serverLoop() {  
    connection = AcceptConnection();  
    ThreadCreate(ServiceWebPage(), connection);  
}
```

- Advantages of threaded version:
 - Can **share file caches** kept in memory, results of PHP scripts, other things
 - Threads are **much cheaper to create** than processes, so this has a lower per-request overhead
- What if too many requests come in at once?

Thread Pools

- Problem with previous version: Unbounded Threads
- Instead, allocate a **bounded “pool” of threads**, representing the maximum level of multiprogramming



```
master() {  
    allocThreads(slave, queue);  
    while(TRUE) {  
        con=AcceptConnection();  
        Enqueue(queue, con);  
        wakeUp(queue);  
    }  
}  
  
slave(queue) {  
    while(TRUE) {  
        con=Dequeue(queue);  
        if (con==null)  
            sleepOn(queue);  
        else  
            ServiceWebPage(con);  
    }  
}
```

Shared states are necessary evil

- Shared states are useful!
 - Shared variables of threads are much cheaper than those of processes
- Shared states are horrible!
 - Programs must be insensitive to arbitrary interleavings. 문이지않.
 - Without careful design, shared variables can output completely inconsistent results

Problem is at the lowest level

- Most of the time, threads are working on separate data, so scheduling doesn't matter:

Thread A

$x = 1;$

Thread B

$y = 2;$

- However, what if on shared data (initially, $y = 0$)?

Thread A

$x = 1;$

$x = y + 1;$

Thread B

$y = 2;$

$y = y * 2;$

- What are the possible values of x ?

Thread A

$x = 1;$

$x = y + 1;$

Thread B

$y = 2;$

$y = y * 2$

$x = 1$

Problem is at the lowest level

- Most of the time, threads are working on separate data, so scheduling doesn't matter:

Thread A

$x = 1;$

Thread B

$y = 2;$

- However, what if on shared data (initially, $y = 0$)?

Thread A

$x = 1;$

$x = y + 1;$

Thread B

$y = 2;$

$y = y * 2;$

- What are the possible values of x ?

Thread A

$x = 1;$

$x = y + 1;$

Thread B

$y = 2;$

$y = y * 2;$

$x = 5$

Problem is at the lowest level

- Most of the time, threads are working on separate data, so scheduling doesn't matter:

Thread A

`x = 1;`

Thread B

`y = 2;`

- However, what if on shared data (initially, $y = 0$)?

Thread A

`x = 1;`

`x = y+1;`

Thread B

`y = 2;`

`y = y*2;`

- What are the possible values of x ?

Thread A

`x = 1;`

`x = y+1;`

Thread B

`y = 2;`

`y = y*2;`

`x=3`

Definitions

- **Race condition**
 - A situation where multiple processes access and manipulate the same data concurrently and the outcome of the execution depends on the particular order (interleaving) of such accesses



More Definitions

- **Critical Section:** a piece of code that accesses a shared resource (producing a race condition)
- **Mutual Exclusion:** ensuring that only one thread executes critical section
 - One thread *excludes* the other(s) while doing its task
- **Lock:** prevent someone from doing something
 - Lock before entering critical section, before accessing shared data
 - Unlock when leaving, after done accessing shared data
 - Wait if locked (all synchronization involves waiting!)

“Too much milk”

- Great thing about OS's – analogy between problems in OS and problems in real life
 - Help you understand real life problems better
- Example: People need to coordinate:



Time	Person A	Person B
3:00	Look in Fridge. Out of milk	
3:05	Leave for store	
3:10	Arrive at store	Look in Fridge. Out of milk
3:15	Buy milk	Leave for store
3:20	Arrive home, put milk away	Arrive at store
3:25		Buy milk
3:30		Arrive home, put milk away

Two guarantees

- **Safety:** A program never enters a bad state
 - Too much milk: Never more than one person buys milk
- **Liveness:** A program eventually a good state
 - Too much milk: If milk is needed, someone eventually buys it

Too Much Milk, Try #1

- Correctness property
 - Someone buys if needed (liveness)
 - At most one person buys (safety)
- Try #1: leave a note

```
if (!milk)
    if (!note) {
        leave note
        buy milk
        remove note
    }
```

Too Much Milk: Solution #1

- Still too much milk **but only occasionally!**

<u>Thread A</u>	<u>Thread B</u>
<pre>if (!Milk) if (!Note) {</pre>	
	<pre>if (!Milk) if (!Note) {</pre>
<pre> leave Note; buy milk; remove note; }</pre>	
<pre>}</pre>	<pre> leave Note; buy milk;</pre>
	<pre> ...</pre>

Thread can get context switched after checking milk and note but before leaving note!

Solution makes problem worse since fails **intermittently**

비행기

Makes it really hard to debug...

Must work despite how threads are interleaved

Too Much Milk, Try #2

Thread A

```
leave note A ①  
if (!note B) { ④  
    if (!milk)  
        buy milk  
}  
remove note A ⑤
```

Thread B

```
② leave note B  
③ if (!noteA) {  
    if (!milk)  
        buy milk  
}  
⑥ remove note B
```

→ Liveness problem

Too Much Milk Solution #2

- Possible for neither thread to buy milk!

Thread A

```
leave note A;
```

Thread B

```
leave note B;  
if (!Note A) {  
    if (!Milk) {  
        buy Milk;  
    }  
}
```


```
if (!Note B) {  
    if (!Milk) {  
        buy Milk;  
    }  
}
```

...

```
remove note B;
```

- Really insidious:
 - **Unlikely** that this would happen, **but possible at worst case**

Too Much Milk Solution #2: problem!



- *I'm* not getting milk, *You're* getting milk
- This kind of lockup is called “**starvation!**”

Too Much Milk, Try #3

Thread A

leave note A

while (note B) // X

do nothing;

if (!milk)

buy milk;

remove note A

Thread B

leave note B

if (!noteA) { // Y

if (!milk)

buy milk

}

remove note B

 Can guarantee at X and Y that either:

- (i) Safe for me to buy
- (ii) Other will buy, ok to quit

Lessons

- Solution is ***complicated***
 - “obvious” code often has bugs
- Generalizing to many threads/processors
 - Even more complex: see Peterson’s algorithm

OS needs to provide a simple way to solve problems like too much milk!

Locks

- Lock::acquire
 - wait until lock is free, then take it
- Lock::release
 - release lock, waking up anyone waiting for it

Formal guarantees

1. At most one lock holder at a time (safety)
2. If no one holding, acquire gets lock (liveness)
3. If all lock holders finish and no higher priority waiters, waiter eventually gets lock (liveness)

Too Much Milk, #4

Locks allow concurrent code to be much simpler:

```
lock.acquire();
```

```
if (!milk)
```

```
    buy milk
```

```
lock.release();
```

Why does the lock require operating system support?

Roadmap

Concurrent Applications

Shared Objects

Bounded Buffer

Barrier

Synchronization Variables

Semaphores

Locks

Condition Variables

Atomic Instructions

Interrupt Disable

Test-and-Set

Hardware

Multiple Processors

Hardware Interrupts

Rules for Using Locks


- Lock is initially free
- Always acquire before accessing shared data structure
 - Beginning of procedure!
- Always release after finishing with shared data
 - End of procedure!
 - Only the lock holder can release
- Never access shared data without lock
 - Danger!

Race condition example

Two threads run on the same bank server

```
withdraw (account, amount) {  
1000 balance = get_balance(account);  
    balance = balance - amount;  
    put_balance(account, balance);  
    return balance;  
}
```

```
withdraw (account, amount) {  
1000 balance = get_balance(account);  
    balance = balance - amount;  
    put_balance(account, balance);  
    return balance;  
}
```



Suppose you have a balance of \$1000

You visit online banking site with two web browsers
and try to withdraw \$100 at the same time

What happens? What is/are possible outcome(s)?

\$900 result balance with 2 x \$100 withdrawals

Lock Example

```
withdraw (account, amount) {  
    acquire(lock);  
    balance = get_balance(account);  
    balance = balance - amount;  
    put_balance(account, balance);  
    release(lock);  
    return balance;  
}
```

**Critical
Section**

```
acquire(lock);  
balance = get_balance(account);  
balance = balance - amount;
```


```
acquire(lock);
```

```
put_balance(account, balance);  
release(lock);
```

```
balance = get_balance(account);  
balance = balance - amount;  
put_balance(account, balance);  
release(lock);
```


Bounded-Buffer (producer-consumer) Problem

- Problem Definition
 - Producer puts things into a shared buffer
 - Consumer takes them out
 - Need **synchronization to coordinate producer/consumer**




- Don't want producer and consumer to have to work in lockstep, so put a fixed-size buffer between them
 - Need to **synchronize access to this buffer**
 - Producer needs to **wait if buffer is full**
 - Consumer needs to **wait if buffer is empty**



Example: Bounded Buffer

Initially: front = tail = 0; lock = FREE; MAX is buffer capacity

```
tryput(item) {  
    lock.acquire();  
    if ((tail - front) < size) {  
        buf[tail % MAX] = item; ②  
        tail++;  
    }  
    lock.release();  
}
```



```
tryget() {  
    item = NULL;  
    lock.acquire();  
    if (front < tail) {  
        item = buf[front % MAX];  
        front++;  
    }  
    lock.release();  
    ③ ↓ return item;  
}
```



Question

- If tryget returns NULL, do we know the buffer is empty?

No

- How can a thread know when a buffer is empty?
 - **Need another primitive for the purpose**

Condition Variables

- Waiting for a change to shared states
 - Called only when holding a lock
- Wait(): atomically release lock and relinquish processor
 - Re-acquire the lock when wakened
- Signal(): wake up a waiter, if any
- Broadcast(): wake up all waiters, if any

Condition Variable Design Pattern

cond_var is a condition variable

```
FunctionThatWaits() {  
    lock.acquire();  
    // Read/write shared state  
  
    while (!testSharedState()) {  
        cond_var.wait(&lock);  
    }  
  
    // Read/write shared state  
    lock.release();  
}
```

```
FunctionThatSignals() {
    lock.acquire();
    // Read/write shared state

    // If testSharedState is now true
    cond_var.signal();

    lock.release();
}
```

만약 $\frac{1}{2} \leq \frac{1}{2} \leq \frac{1}{2}$

Conditional variable is memoryless

- CV does not have internal states other than a queue of waiting thread
- If no threads are in the waiting queue, a signal or broadcast has no effect
- CV does not have memory of earlier calls of signal or broadcast

Does it work?

```
methodThatWaits() {  
    lock.acquire();  
    // Read/write shared state  
    lock.release();
```

```
    lock.acquire();  
    while (!testSharedState()) {  
        cv.wait(&lock);  
    }  
    lock.release();
```

```
    lock.acquire();  
    // Read/write shared state  
    lock.release();  
}
```

```
methodThatSignals() {  
    lock.acquire();  
    // Read/write shared state  
  
    // If testSharedState is now true  
    cv.signal();  
  
    lock.release();  
}
```

여기서 lock을 해제함

→ 조건을 확인함.

Bounded buffer

- Producer

```
int i, loop = MAX_LOOP;  
For (i = 0; i < loop; i++) {  
    put(i)  
}
```

- Consumer

```
int tmp;  
while(1) {  
    tmp = get();  
    printf("%d\n", tmp);  
}
```



Safety problem

Try 1: Bounded Buffer

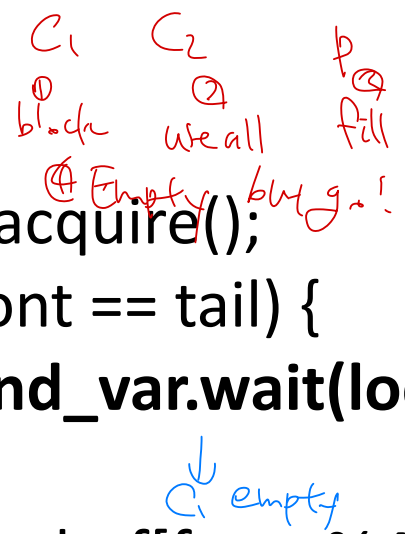


Initially: front = tail = 0; MAX is buffer capacity
cond_var are condition variables

```
put(item) {  
    lock.acquire();  
    if ((tail - front) == MAX) {  
        cond_var.wait(lock);  
    }  
    ① buf[tail % MAX] = item;  
    tail++;  
    cond_var.signal(lock);  
    lock.release();  
}
```



```
get() {  
    lock.acquire();  
    if (front == tail) {  
        cond_var.wait(lock);  
    }  
    item = buf[front % MAX];  
    front++;  
    cond_var.signal(lock);  
    lock.release();  
    return item;  
}
```



Try 2: Bounded Buffer

Initially: front = tail = 0; MAX is buffer capacity
cond_var are condition variables

```

put(item) {
    lock.acquire();
    while ((tail - front) == MAX) {
        cond_var.wait(lock);
    }
    buf[tail % MAX] = item;
    tail++;
    cond_var.signal(lock);
    lock.release();
}
    
```

②

```

get() {
    lock.acquire();
    while (front == tail) {
        cond_var.wait(lock);
    }
    item = buf[front % MAX];
    front++;
    cond_var.signal(lock);
    lock.release();
    return item;
}
    
```

① ②
⑤ → empty
④ ②
③

Solution: Bounded Buffer

Initially: front = tail = 0; MAX is buffer capacity
empty/full are condition variables

The diagram illustrates a swap in the implementation of the bounded buffer. A red 'X' is drawn across the two code blocks, indicating that the **full** and **empty** condition variables are swapped. A blue arrow points from the **full** condition variable in the **put** function to the **empty** condition variable in the **get** function. Another blue arrow points from the **empty** condition variable in the **put** function to the **full** condition variable in the **get** function.

```
put(item) {  
    lock.acquire();  
    while ((tail - front) == MAX) {  
        full.wait(lock);  
    }  
    buf[tail % MAX] = item;  
    tail++;  
    empty.signal(lock);  
    lock.release();  
}  
  
get() {  
    lock.acquire();  
    while (front == tail) {  
        empty.wait(lock);  
    }  
    item = buf[front % MAX];  
    front++;  
    full.signal(lock);  
    lock.release();  
    return item;  
}
```

Summary: Condition Variables

- ALWAYS hold lock when calling wait, signal, broadcast
 - Condition variable is sync FOR shared state
 - ALWAYS hold lock when accessing shared state
- Condition variable is memoryless
 - If signal when no one is waiting, no op
 - If wait before signal, waiter wakes up

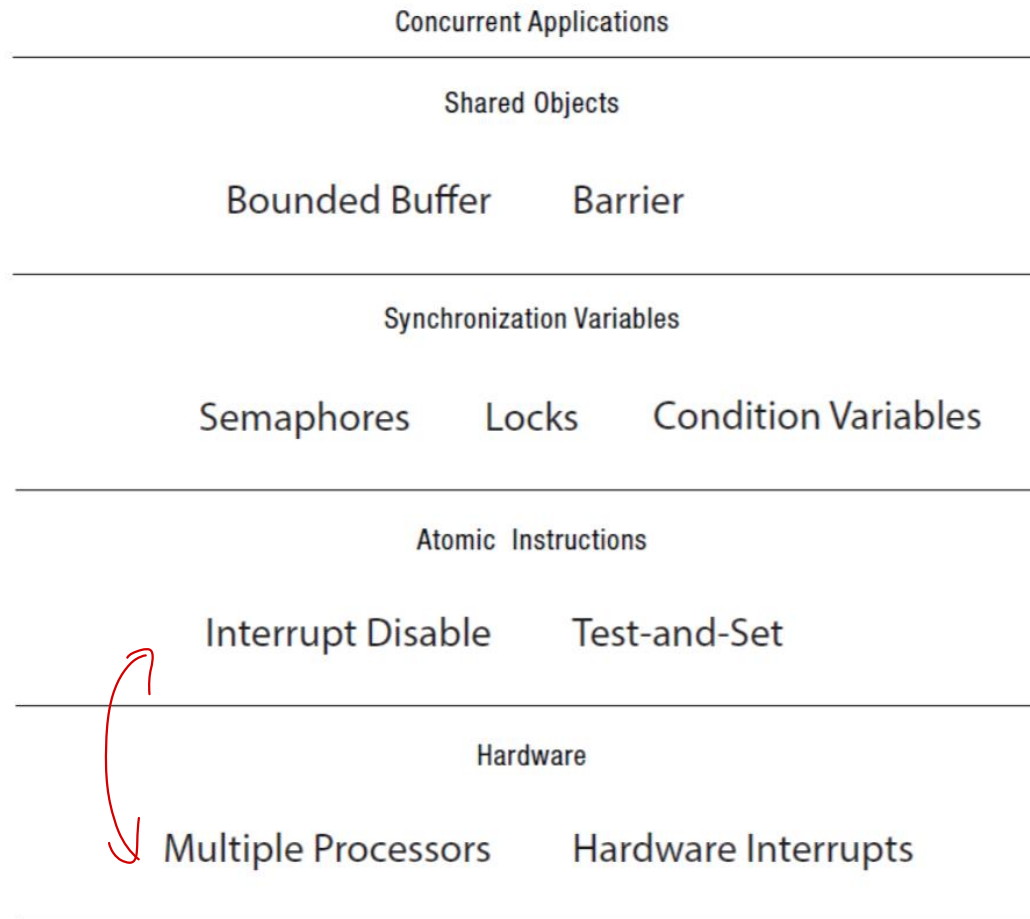
Summary: Condition Variables, cont'd

- When a thread is woken up from wait, it may not run immediately
 - Signal/broadcast put thread on ready list
 - When lock is released, anyone might acquire it
- Wait MUST be in a loop
 - while (needToWait()) {
 - condition.Wait(lock);
 - }

Remember the rules

- Always use locks and condition variables for shared states *Memoryless!*
- Always acquire lock at beginning of procedure, release at end
- Always hold lock when using a condition variable
- Always wait in while loop
- Never use sleep() to wait for a thread to finish a task

Implementing Synchronization



How to Implement Lock?

- **Lock:** prevents someone from accessing something
 - Lock before entering critical section (e.g., before accessing shared data)
 - Unlock when leaving, after accessing shared data
 - Wait if locked
 - **Important idea: all synchronization involves waiting**
 - **Should sleep if waiting for long time**



Naïve use of Interrupt Enable/Disable

- How can we build ["]atomic operations["]?
 - Recall: A thread loses its control in two ways.
 - (Internal: Relinquishing the CPU (e.g., sleep, yield)
 - External: Interrupts (ex. timer)
- On a uniprocessor, can avoid context-switching by:
 - Avoiding internal events
 - Preventing external events by disabling interrupts

Consequently, naïve Implementation of locks:

```
LockAcquire { disable Ints; }
```

```
LockRelease { enable Ints; }
```

(Why implemented by kernel? : These are privileged instructions

Disabling interrupt

- Privileged instruction or not?

Yes

- If code between enabling/disabling interrupt is long, then what is a problem?

Concurrency: For example, if timer interrupt got off, then every other system with timer gets lost (ex, scheduler, ...)

Lock Implementation, Uniprocessor

```
Int lockValue = FREE;
```

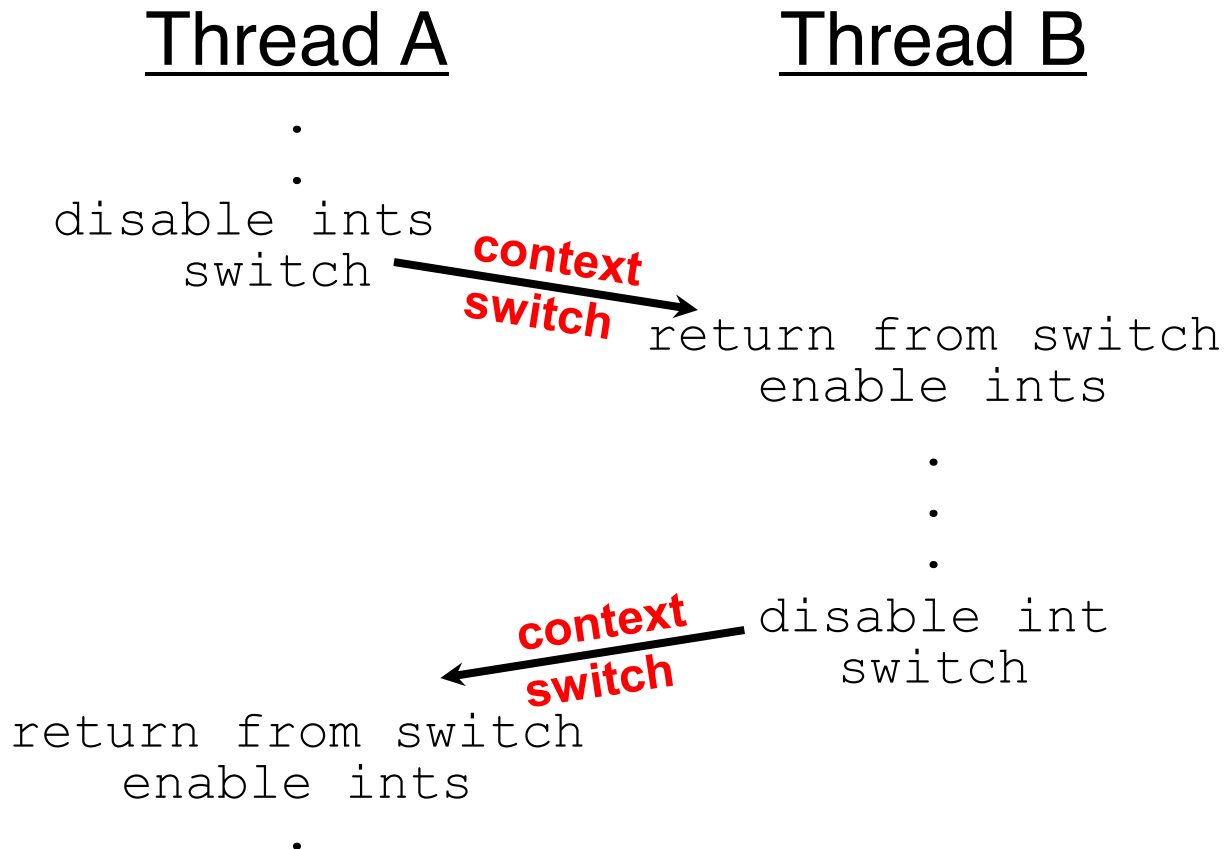
```
Lock::acquire() {  
    disableInterrupts();  
    if (lockValue == BUSY) {  
        waiting.add(TCB);  
        TCB->state = WAITING;  
        next = readyList.remove();  
        switch(TCB, next);  
        TCB->state = RUNNING;  
    } else {  
        lockValue = BUSY;  
    }  
    enableInterrupts();  
}
```

← Code goes here

```
Lock::release() {  
    disableInterrupts();  
    if (!waiting.Empty()) {  
        next = waiting.remove();  
        next->state = READY;  
        readyList.add(next);  
    } else {  
        lockValue = FREE;  
    }  
    enableInterrupts();  
}
```

How to Re-enable After switch()?

- Since ints are disabled when you call sleep:
 - Responsibility of the **next thread to re-enable ints**
 - When the sleeping thread wakes up, returns to acquire and re-enables interrupts



Enable/disable interrupt in multiprocessors

- Does it guarantee atomic instructions like uniprocessor case?

No,

E/D interrupt is only per processor.
Disable
→ [S] CPU1 [S] CPU2 ← Already running

What can OS do to provide atomic instructions?

Using global variable.

Spinlocks with busy loop

A spinlock is a lock where the processor waits in a loop for the lock to become free


- Assumes lock will be held for a short time
- Used to protect the CPU scheduler and to implement locks

```
Spinlock::acquire() {  
    while (lockValue == BUSY)  
        ;  
    lockValue = Busy  
}
```

```
Spinlock::release() {  
    lockValue = FREE;  
}
```

Problems in multiprocessors?

→ HW help



Multiprocessor

- Atomic read-modify-write instructions
 - Atomically read a value from memory, operate on it, and then write it back to memory
 - Intervening instructions prevented in hardware
- Examples 예시
 - Test and set
 - Intel: xchgb, lock prefix
 - Compare and swap
- Any of these can be used for implementing locks and condition variables!

Spinlock with test-and-set instruction

```
BUSY = 1; FREE = 0;
```

```
Int lockValue = FREE;
```

What CPU does atomically:

```
int TestAndSet(int *old_ptr) {  
    int old = *old_ptr; // fetch old value at old_ptr  
    *old_ptr = BUSY;    // store BUSY into old_ptr  
    return old;         // return the old value  
}
```

```
Spinlock::acquire() {  
    while(testAndSet(&lockValue)  
        == BUSY)  
        ;  
}
```

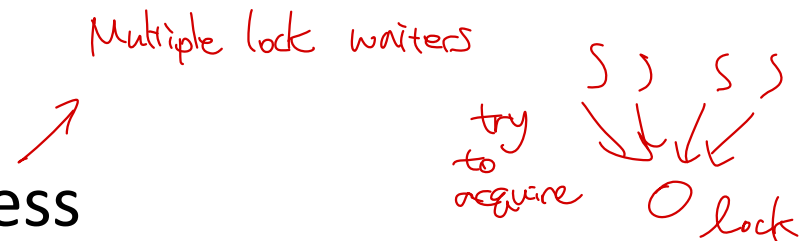
```
Spinlock::release() {  
    lockValue = FREE;  
}
```


Two main problems in the previous spinlock implementation

- It is spinning!
 - a thread waiting for the spinlock waste CPU cycles
 - Especially severe problem in uniprocessor

- Does not guarantee liveness

- A thread may keep trying to hold the spinlock



Avoid spinning

```
int TestAndSet(int *old_ptr) {  
    int old = *old_ptr; // fetch old value at old_ptr  
    *old_ptr = BUSY;    // store BUSY into old_ptr  
    return old;         // return the old value  
}
```

```
Spinlock::acquire() {  
    while(testAndSet(&lockValue)  
        == BUSY)  
        yield();  
}
```

yield CPU

```
Spinlock::release()  
{  
    lockValue = FREE;  
}
```

New HW primitive: fetch-and-add

also atomic

- ```
int FetchAndAdd(int *ptr) {
 int old = *ptr;
 *ptr = old + 1;
 return old;
}
```
- Any ideas to solve the starvation of the previous spinlock implementation?

# Ticket locks

Global variable

```
typedef struct __lock_t {
 int ticket;
 int turn;
} lock_t;
```

```
void lock_init(lock_t *lock){
 lock->ticket = 0;
 lock->turn = 0;
}
```

```
void lock(lock_t *lock) {
 int myturn = FetchAndAdd(&lock->ticket);
 while (lock->turn != myturn)
 ; // spin
}
```

```
void unlock(lock_t *lock) {
 lock->turn = lock->turn + 1;
}
```

# Semaphores



- Semaphores are a kind of generalized locks
  - First defined by Dijkstra in late 60s
  - Main synchronization primitive used in original UNIX
- Definition: a Semaphore has a non-negative integer value and supports the following two operations:
  - **P()**: an atomic operation that waits for semaphore to become positive, then decreases it by 1
    - Think of this as the **wait()** operation
  - **V()**: an atomic operation that increases the semaphore by 1, waking up a waiting P, if any
    - Think of this as the **signal()** operation
  - Note that P() stands for “*proberen*” (to test) and V() stands for “*verhogen*” (to increment) in Dutch

# Two Uses of Semaphores (1)

- **Mutual Exclusion** (initial value = 1)


- Also called “Binary Semaphore”.
- Can be used for mutual exclusion:

```
semaphore.P(); //wait
// Critical section goes here
semaphore.V(); //signal
```

- **Counting semaphore** – integer value can range over an unrestricted domain

# Semaphores

- Semaphore from railway analogy
  - Here is a semaphore initialized to 2 for resource control:



# Two Uses of Semaphores (2)

- Scheduling Constraints (initial value = 0)
  - Allow thread 1 to wait for a signal from thread 2, i.e., thread 2 **schedules** thread 1 when a given **constraint** is satisfied
  - Example: suppose you had to implement ThreadJoin which must wait for the thread to terminate:

Initial value of semaphore = 0

```
ThreadJoin {
 semaphore.P();
}
```

```
ThreadFinish {
 semaphore.V();
}
```





# Semaphores

- Semaphore from railway analogy
  - Here is a semaphore initialized to 0 for scheduling constraint:
    - One train leaves only after another train comes



---

**Value=0**

---

# Semaphore Bounded Buffer

Initially: front = last = 0; MAX is buffer capacity  
mutex = 1; emptySlots = MAX; fullSlots = 0;

```
get() {
 fullSlots.P();
 mutex.P();
 item = buf[front % MAX];
 front++;
 mutex.V();
 emptySlots.V();
 return item;
}
```

```
put(item) {
 emptySlots.P();
 mutex.P();
 buf[last % MAX] = item;
 last++;
 mutex.V();
 fullSlots.V();
}
```

# Recall: conditional variable example

Initially: front = tail = 0; MAX is buffer capacity  
empty/full are condition variables

The diagram illustrates a swap in the implementation of the conditional variable example. Two functions, `put(item)` and `get()`, are shown side-by-side. In the original code, `put()` calls `full.wait(lock)` and `empty.signal()`, while `get()` calls `empty.wait(lock)` and `full.signal()`. Red lines and circles indicate a swap: the `while` loops in both functions are circled in red, and red lines cross the `wait` and `signal` calls between the two functions. This results in the modified code shown below.

```
put(item) {
 lock.acquire();
 while ((tail - front) == MAX) {
 full.wait(lock);
 }
 buf[tail % MAX] = item;
 tail++;
 empty.signal();
 lock.release();
}

get() {
 lock.acquire();
 while (front == tail) {
 empty.wait(lock);
 }
 item = buf[front % MAX];
 front++;
 full.signal();
 lock.release();
 return item;
}
```

# Implementing Condition Variables using Semaphores (Try 1)

```
wait(lock) {
 lock.release();
 semaphore.P();
 lock.acquire();
}

signal() {
 semaphore.V();
}
```

Conditional Variable

CV = memoryless

⚡ Semaphore has memory

Solution: Only signal when  
wait is called

# Implementing Condition Variables using Semaphores (Try 2)

```
wait(lock) {
```

```
 ① lock.release();
```

```
 ② semaphore.P();
```

```
 lock.acquire();
}
```

→ 3 in one is not atomic

Signal lost!

```
signal() {
```

```
 ③ if (semaphore is not empty)
```

```
 semaphore.V();
```

→ Not called when empty

# Implementing Condition Variables using Semaphores

```
wait(lock) {
 semaphore = new Semaphore;
 queue.Append(semaphore); // queue of waiting threads
 lock.release();
 semaphore.P();
 lock.acquire();
}
signal() {
 if (!queue.Empty()) {
 semaphore = queue.Remove();
 semaphore.V(); // wake up waiter
 }
}
```

①

②

Wait skip

# Remember the rules

- Use consistent structure
- Always use locks and condition variables
- If you are not sure, always acquire lock at beginning of procedure, release at end
- Always hold lock when using a condition variable
- Always wait in while loop